

# A Cryptographic Analysis of the TLS 1.3 Handshake Protocol

Benjamin Dowling<sup>1</sup>, Marc Fischlin<sup>2</sup>, Felix Günther<sup>1</sup>, and Douglas Stebila<sup>3</sup>

<sup>1</sup>Department of Computer Science, ETH Zürich

<sup>2</sup>TU Darmstadt

<sup>3</sup>University of Waterloo

February 22, 2021

## Abstract

We analyze the handshake protocol of the Transport Layer Security (TLS) protocol, version 1.3. We address both the full TLS 1.3 handshake (the one round-trip time mode, with signatures for authentication and (elliptic curve) Diffie–Hellman ephemeral ((EC)DHE) key exchange), and the abbreviated resumption/“PSK” mode which uses a pre-shared key for authentication (with optional (EC)DHE key exchange and zero round-trip time key establishment). Our analysis in the reductionist security framework uses a multi-stage key exchange security model, where each of the many session keys derived in a single TLS 1.3 handshake is tagged with various properties (such as unauthenticated versus unilaterally authenticated versus mutually authenticated, whether it is intended to provide forward security, how it is used in the protocol, and whether the key is protected against replay attacks). We show that these TLS 1.3 handshake protocol modes establish session keys with their desired security properties under standard cryptographic assumptions.

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Development and Standardization of TLS 1.3 . . . . .	3
1.2	Security Analyses of TLS . . . . .	4
1.3	Our Contributions . . . . .	4
<b>2</b>	<b>Preliminaries</b>	<b>7</b>
2.1	Notation . . . . .	7
2.2	Collision-Resistant Hash Functions . . . . .	7
2.3	HMAC and HKDF . . . . .	8
2.4	Dual PRF Security and the PRF-ODH Assumption . . . . .	8
<b>3</b>	<b>The TLS 1.3 Handshake Protocol</b>	<b>9</b>
3.1	Key-Exchange Phase . . . . .	9
3.2	Authentication Phase . . . . .	11
3.3	NewSessionTicket . . . . .	13
<b>4</b>	<b>Multi-Stage Key Exchange Security Model</b>	<b>14</b>
4.1	Syntax . . . . .	17
4.2	Adversary Model . . . . .	20
4.3	Security of Multi-Stage Key Exchange Protocols . . . . .	22
4.3.1	Match Security . . . . .	22
4.3.2	Multi-Stage Security . . . . .	24
<b>5</b>	<b>Security Analysis of the TLS 1.3 Full 1-RTT Handshake</b>	<b>25</b>
5.1	Match Security . . . . .	26
5.2	Multi-Stage Security . . . . .	27
<b>6</b>	<b>Security Analysis of the TLS 1.3 PSK/PSK-(EC)DHE (with Optional 0-RTT) Handshakes</b>	<b>32</b>
6.1	TLS 1.3 PSK-only (0-RTT optional) . . . . .	33
6.1.1	Match Security . . . . .	33
6.1.2	Multi-Stage Security . . . . .	35
6.2	TLS 1.3 PSK-(EC)DHE (0-RTT optional) . . . . .	38
6.2.1	Match Security . . . . .	38
6.2.2	Multi-Stage Security . . . . .	40
<b>7</b>	<b>Discussion and Conclusions</b>	<b>48</b>
7.1	Technical Differences from Our Earlier Work . . . . .	48
7.2	Comments on the TLS 1.3 Design . . . . .	49
7.3	Open Research Questions . . . . .	52
7.4	Conclusions . . . . .	53
<b>A</b>	<b>Reducing Multiple to Single Test Queries</b>	<b>61</b>

# 1 Introduction

The *Transport Layer Security (TLS)* protocol is one of the most widely deployed cryptographic protocols in practice, protecting numerous web and e-mail accesses every day. The *TLS handshake protocol* allows a client and a server to authenticate each other and to establish a key, and the subsequent *record layer protocol* provides confidentiality and integrity for communication of application data. Originally developed as the Secure Sockets Layer (SSL) protocol version 3 in 1996, TLS version 1.0 was standardized by the Internet Engineering Task Force (IETF) in 1998 [DA99], with subsequent revisions to version 1.1 (2006) [DR06] and version 1.2 (2008) [DR08]. Despite its large-scale deployment, or perhaps because of it, we have witnessed frequent successful attacks against TLS. Starting around 2009, there were many practical attacks on the then-current version 1.2 of TLS that received significant attention, exploiting weaknesses in underlying cryptographic primitives (such as weaknesses in RC4 [ABP<sup>+</sup>13]), errors in the design of the TLS protocol (e.g., BEAST [Duo11], the Lucky 13 attack [AP13], the triple handshake attack [BDF<sup>+</sup>14], the POODLE attack [MDK14], the Logjam attack [ABD<sup>+</sup>15]), or flaws in implementations (e.g., the Heartbleed attack [Cod14], state machine attacks (SMACK [BBDL<sup>+</sup>15])).

## 1.1 Development and Standardization of TLS 1.3

With concerns rising about the security of TLS version 1.2 due to the many attacks, but also motivated by desire to deprecate old algorithms, enhance privacy, and reduce connection establishment latency, in 2014 the IETF’s TLS working group initiated a multi-year process to develop and standardize a new version of TLS, eventually called version 1.3. From 2014 through 2018, a total 29 drafts of TLS 1.3 were published, with active feedback from industry and academia, including extensive security analyses by various teams from academia (see [PvdM16] for a chronicle of the development and analysis of TLS 1.3). The document standardizing TLS 1.3, RFC 8446 [Res18], was published in August 2018 and has already seen widespread adoption.

From a cryptographic perspective, major design changes in TLS 1.3 compared to version 1.2 include: (1) encrypting some handshake messages with an intermediate session key, to provide confidentiality of handshake data such as the client certificate; (2) signing the entire handshake transcript for authentication; (3) including hashes of handshake messages in a variety of key calculations; (4) using different keys to encrypt handshake messages and application data; (5) deprecating a variety of cryptographic algorithms (including RSA key transport, finite-field Diffie–Hellman key exchange, SHA-1, RC4, CBC mode, MAC-then-encode-then-encrypt); (6) using modern authenticated encryption with associated data (AEAD) schemes for protecting application data; and (7) providing handshakes with fewer message flows to reduce latency.

There are two primary modes of the TLS 1.3 handshake protocol. One is the full, one round-trip time (1-RTT) handshake, which uses public-key certificates for server and (optionally) client authentication, and (elliptic curve) Diffie–Hellman ephemeral ((EC)DHE) key exchange, inspired by Krawczyk’s ‘SIGn-and-Mac’ (SIGMA) design [Kra03]. Several session keys are established for a variety of purposes in this mode: to encrypt part of the handshake, to enable export of keying material to other applications, for session resumption, and of course to encrypt application data. This mode gets its name from the fact that application data can be sent from the client to the server with the handshake’s completion after a full round trip, meaning there is one round-trip time (1-RTT) until the first application message can be sent (not counting non-TLS networking operations such as DNS lookups or the TCP 3-way handshake).

The other primary mode of the TLS 1.3 handshake protocol is the resumption or pre-shared key (PSK) mode, in which authentication is based on a symmetric pre-shared key, with optional

(EC)DHE key exchange for forward secrecy; this generalizes the abbreviated session resumption handshake from earlier versions of TLS. The PSK mode can optionally be augmented with a zero round-trip time (0-RTT) key establishment, allowing the client to send—along with its first TLS flow—application data encrypted under a key derived from the PSK.

## 1.2 Security Analyses of TLS

**TLS 1.2 and prior versions.** A long line of work has analyzed various versions of the SSL/TLS protocol using both formal methods and reductionist security proofs. In the reductionist security paradigm, early work [JK02, MSW08, Gaj08] on the handshake protocol dealt with modified or truncated versions of the protocol, necessary because TLS 1.2 and earlier did not have strict key separation: the session key was also used to encrypt messages within the handshake protocol, barring security proofs in strong indistinguishability-based authenticated key exchange models in the Bellare–Rogaway [BR94] style. There were also formalizations of the security of the authenticated encryption in the record layer [Kra01, PRS11]. A major milestone in reductionist analyses of TLS was the development of the authenticated and confidential channel establishment (ACCE) security model which allowed for the combined analysis of a full TLS 1.2 handshake and secure channel in a single model [JKSS12], sidestepping the aforementioned key separation issue; this work was followed by a range of other works analyzing the security of various aspects of TLS 1.2 [KPW13, KSS13, LSY<sup>+</sup>14, GKS13, DS15]. Other approaches to proving the security of TLS 1.2 within the reductionist security paradigm include a range of modular and compositional approaches [BFS<sup>+</sup>13] as well as approaches that combine formal analysis and reductionist security [BFK<sup>+</sup>13, BFK<sup>+</sup>14].

**TLS 1.3 drafts.** The handshake protocol in initial drafts of TLS 1.3 was based in part on the OPTLS protocol [KW16]. There were a variety of investigations on the security of various drafts throughout the TLS 1.3 standardization process. Using the reductionist security paradigm, there have been analyses of the handshake protocol [DFGS15, KMO<sup>+</sup>15, DFGS16, KW16, LXZ<sup>+</sup>16, FGSW16, Kra16b, FG17, BFG19a] and the record layer [BMM<sup>+</sup>15, BT16, LP17, GM17, PS18]. There has been a range of work involving formal methods and tools, such as model checkers and symbolic analysis [CHSv16, CHH<sup>+</sup>17], and approaches combining verified implementations with formal analysis and reductionist security [BFK16, BBD<sup>+</sup>15, BBF<sup>+</sup>16, DFK<sup>+</sup>17].

**TLS 1.3 standard.** Since TLS 1.3 was published as an RFC in August 2018, some works have addressed the final TLS 1.3 standard. The Selfie attack [DG21b] led to updated analyses of PSK handshakes [DG21b, AASS19]. Arfaoui et al. [ABF<sup>+</sup>19] investigated the privacy features of the TLS 1.3 handshake. Revised computational security proofs of the full 1-RTT handshake by Diemert and Jager [DJ21] and Davis and Günther [DG21a] translated techniques of Cohn-Gordon et al. [CCG<sup>+</sup>19] to establish tighter reductions. There have also been academic proposals for improvements to or modifications of TLS 1.3, considering forward security for the 0-RTT handshake [AGJ21], running TLS 1.3 over a different network protocol [CJJ<sup>+</sup>19], or defining a KEM-based alternative handshake enabling the deployment of post-quantum schemes [SSW20].

## 1.3 Our Contributions

We give a reductionist security analysis of three modes of the TLS 1.3 handshake: the full 1-RTT handshake, the PSK handshake (with optional 0-RTT mode), and the PSK-(EC)DHE handshake (with optional 0-RTT mode); based on a cryptographic abstraction of the protocols we provide

in Section 3. In order to carry out our analysis, we formalize a multi-stage key exchange security model which can capture a variety of characteristics associated to each stage key. Our analysis shows that the design of the TLS 1.3 handshake follows sound cryptographic principles.

**Security model.** Our security model, given in Section 4, follows the Bellare–Rogaway (BR) model [BR94] for authenticated key exchange security based on session key indistinguishability, as formalized by Brzuska et al. [BFWW11, Brz13], and our model builds specifically on the multi-stage model of Fischlin and Günther [FG14, Gün18]. The latter deals with key exchange protocols that derive a series of session keys in the course of multiple protocol stages. Our extension of their multi-stage key exchange model allows us to capture the following characteristics associated to the session key established at each stage, which we call the stage key:

- *Authentication*: whether a stage key is unauthenticated, unilaterally authenticated, or mutually authenticated. We further extend the multi-stage model to capture *upgradable authentication*: a stage’s key may be considered, say, unauthenticated at the time it is accepted, but the authentication level of this key may be “raised” to unilaterally authenticated or, potentially in a second step, mutually authenticated after some later operations, such as verification of a signature in a later message.
- *Forward secrecy*: whether a stage key is meant to provide forward secrecy, namely that it remains secure after compromise of a long-term secret involved in its derivation.
- *Key usage*: whether a stage key is meant to be used internally within the protocol (for example, to encrypt later handshake messages), or externally (for example, composed with a symmetric encryption scheme to protect application messages or used in some other external symmetric-key protocol).
- *Replayability*: whether it is guaranteed that a stage key is not established in result of a replay attack; early stages of the 0-RTT modes do not have this guarantee.

Our security model comes in two flavors that capture security established through two types of credentials: public keys or symmetric pre-shared keys. Following the BR model, our model of compromise includes long-term key compromise (**Corrupt**) and stage key compromise (**Reveal**). While other models [CK01, LLM07] further capture the compromise of session state or ephemeral randomness, TLS is not designed to be secure against such exposure of ephemeral values and we hence do not include these compromise capabilities in our model.

In addition to capturing indistinguishability of stage keys, the model also ensures soundness of session identifiers using the **Match**-security notion of [BFWW11, Brz13].

**Protocol analysis.** We apply our multi-stage key exchange security model in Sections 5 and 6 to analyze the three modes of the TLS 1.3 handshake: full 1-RTT, PSK, and PSK-(EC)DHE, with the latter two having optional 0-RTT keys. There are four main classes of stage keys covered in the analysis: early data encryption and export keys (ETS, EEMS, only present in the PSK with 0-RTT modes); handshake traffic secrets ( $\text{tk}_{\text{chs}}$ ,  $\text{tk}_{\text{shs}}$ ); application traffic secrets (CATS, SATS); and exported keys (RMS for session resumption, EMS for other exported keys). This results in six stage keys in the full 1-RTT mode and eight stage keys in the PSK modes.

As noted above, our security model allows us to precisely capture various characteristics of different stage keys. For example, consider the client handshake traffic secret  $\text{tk}_{\text{chs}}$ , used to encrypt handshake messages from the client to the server. In the full 1-RTT handshake, this key is initially

unauthenticated, then unilaterally authenticated through a server signature after stage 3 is reached, and may ultimately be mutually authenticated after stage 6 is reached if the client authenticates; it is forward secret; is intended for internal use within the protocol; and it is guaranteed to be non-replayed. In contrast, in the PSK handshake, this key is mutually authenticated as soon as it is established, but does not have forward secrecy. Finally, in the PSK-EC(DHE) handshake, this key is unauthenticated initially, then is upgraded to unilateral and eventually mutual authentication after stages 5 and 8, when MACs within the `Finished` messages are verified; and it is forward secret.

The reductions showing the security of the protocol modes in the model follow a game hopping technique, and mainly rely on standard signature resp. MAC scheme unforgeability (for authentication in the full 1-RTT resp. PSK handshake), hash function collision resistance, PRF security (and in some cases dual PRF security), and an interactive Diffie–Hellman assumption (a variant of the PRF-Oracle-Diffie–Hellman assumption [JKSS12, BFGJ17] called `dual-snPRF-ODH`).

**Observations on the design and security of TLS 1.3.** In Section 7, we include a discussion about various characteristics of TLS 1.3 based on results of our security analysis, including how a variety of TLS 1.3 design decisions positively impact the security analysis (key separation and key independence, including the session hash in signatures and key derivation), some subtleties on the role of handshake encryption and key confirmation via `Finished` messages, as well as the susceptibility of 0-RTT keys to replays.

**Relation to our earlier work.** This paper is successor work to [DFGS15, DFGS16] and [FG17], as well as [Dow17, Gün18]. In [DFGS15], we first extended the multi-stage key exchange model of [FG14] as needed, then applied it to analyze two early drafts of TLS 1.3: `draft-05`, which has the same basic signed-Diffie–Hellman structure but a simplified key schedule compared to the final version, and an alternative proposal called `draft-dh` incorporating ideas from the OPTLS design [KW16], in which servers could have a semi-static DH key share. In [DFGS16], we updated our analysis to `draft-10` and added an analysis of the, by then revised, pre-shared-key handshake mode. In [FG17], a subset of us analyzed the 0-RTT pre-shared key and PSK-(EC)DHE mode in `draft-14`, as well as the later deprecated Diffie–Hellman-based 0-RTT mode using semi-static DH key shares in `draft-12`, which introduced the notion of replayable stages into the multi-stage key exchange security model. In a PhD thesis [Dow17], one of us updated the work from [DFGS16] to address the full, PSK, and PSK-(EC)DHE handshakes in `draft-16`; in another PhD thesis [Gün18], another of us unified the MSKE model and the aforementioned results on the full and PSK handshakes of `draft-10` and the 0-RTT handshakes of `draft-12` and `draft-14`.

This paper updates this prior work to the final version of TLS 1.3 as published in RFC 8446 [Res18] (recall that there were 29 drafts leading up to the final standard). It addresses, in a unified security model, the full, PSK, and PSK-(EC)DHE handshakes, the latter two with optional 0-RTT keys. The security model in this paper includes enhancements not present in earlier works, particularly for capturing upgradable authentication. The model and analysis for the PSK mode have been updated to reflect the observations of Drucker and Gueron’s “Selfie” attack [DG21b] by associating intended roles with a pre-shared key.

Section 7.1 provides more details on technical differences between this paper and our earlier work.

**Limitations.** The TLS 1.3 protocol allows users to support and negotiate different cryptographic algorithms including the used signature schemes, Diffie–Hellman groups, and authenticated en-

encryption schemes. Many implementations will simultaneously support TLS 1.3, TLS 1.2, and even earlier versions. We do not aim to capture the security of this negotiation process nor security when a cryptographic key (e.g., a signing key) is re-used across different algorithm combinations or with earlier versions of TLS [JSS15]. For the PSK modes of TLS 1.3, we do not treat how parties negotiate which pre-shared key to use. Our analysis assumes that all parties use only TLS 1.3 with a single combination of cryptographic algorithms and do not re-use keying material outside of that context (beyond consuming session keys established by the TLS 1.3 handshake).

In our proofs of key indistinguishability for all three TLS 1.3 handshake modes, some of our proof steps involve guessing parties and/or sessions, and thus are non-tight, similar to most proofs of authenticated key exchange protocols. Recently, Diemert and Jager [DJ21] as well as Davis and Günther [DG21a] have established new security proofs for the TLS 1.3 full 1-RTT handshake with tight reductions to the strong Diffie–Hellman assumption, translating techniques of Cohn-Gordon et al. [CCG<sup>+</sup>19].

Our focus is entirely on the TLS 1.3 handshake protocol, and thus does not address security of the record layer’s authenticated encryption. TLS 1.3 also includes a variety of additional functionalities outside the core handshake that we treat as out of scope. Examples include session tickets, post-handshake authentication [Kra16b], the alert protocol, and changes for Datagram TLS (DTLS) 1.3 [RTM19], as well as other extensions to TLS 1.3 currently in the Internet-Draft state.

Security in practice obviously relies on many more factors, such as good implementations and good operational security, which are important but outside the scope of this analysis.

## 2 Preliminaries

We begin with introducing the basic notation we use in this paper and recapping some core building blocks and cryptographic assumptions employed in our security analysis.

### 2.1 Notation

With  $\mathbb{N}$  we denote the natural numbers. We write a bit as  $b \in \{0, 1\}$  and a (bit) string as  $s \in \{0, 1\}^*$ , with  $|s|$  indicating its (binary) length;  $\{0, 1\}^n$  is the set of bit strings of length  $n$ . We write  $x \leftarrow y$  for the assignment of value  $y$  to the variable  $x$  and  $x \leftarrow_{\$} X$  for uniformly sampling  $x$  from a (finite) set  $X$ .

For an algorithm  $\mathcal{A}$  we write  $x \leftarrow \mathcal{A}(y)$ , resp.  $x \leftarrow_{\$} \mathcal{A}(y)$ , for the algorithm deterministically, resp. probabilistically, outputting  $x$  on input  $y$ . We indicate by  $\mathcal{A}^{\mathcal{O}}$  an algorithm  $\mathcal{A}$  running with oracle access to some other algorithm  $\mathcal{O}$ .

### 2.2 Collision-Resistant Hash Functions

As often the case in practice, the cryptographic hash functions used in TLS 1.3 are unkeyed. When considering a hash function’s collision resistance, we hence demand that a security reduction provides effective means for constructing a concrete algorithm generating a collision (cf. Rogaway [Rog06]).

**Definition 2.1** (Hash function and collision resistance). *A hash function  $H: \{0, 1\}^* \rightarrow \{0, 1\}^\lambda$  maps arbitrary-length messages  $m \in \{0, 1\}^*$  to a hash value  $H(m) \in \{0, 1\}^\lambda$  of fixed length  $\lambda \in \mathbb{N}$ .*

*We can now measure the collision resistance (COLL) with respect to an adversary  $\mathcal{A}$  via the advantage*

$$\text{Adv}_{H, \mathcal{A}}^{\text{COLL}} := \Pr [(m, m') \leftarrow_{\$} \mathcal{A} : m \neq m' \text{ and } H(m) = H(m')].$$



In the common asymptotic notion we would demand that one cannot construct an efficient adversary  $\mathcal{A}$  where this advantage is non-negligible with respect to the security parameter  $\lambda$ .

### 2.3 HMAC and HKDF

TLS 1.3 employs HKDF [Kra10, KE10] as its key derivation function, with HMAC [BCK96, KBC97] at its core. We briefly recap their definition and usage.

HMAC [BCK96, KBC97] is based on a cryptographic hash function  $H: \{0, 1\}^* \rightarrow \{0, 1\}^\lambda$  and keyed with some key  $K \in \{0, 1\}^\lambda$  (larger key material is hashed through  $H$  to obtain a  $\lambda$ -bit key). Computing the HMAC value on some message  $m$  is then defined as  $\text{HMAC}(K, m) := H((K \oplus \text{opad}) \parallel H((K \oplus \text{ipad}) \parallel m))$ , where  $\text{opad}$  and  $\text{ipad}$  are two  $\lambda$ -bit padding values consisting of repeated bytes  $0x5c$  and  $0x36$ , respectively.

HKDF follows the *extract-then-expand* paradigm for key derivation [Kra10, KE10], instantiated with HMAC. We adopt the standard notation for the two HKDF functions:  $\text{HKDF.Extract}(XTS, SKM)$  on input an (non-secret and potentially fixed) extractor salt  $XTS$  and some (not necessarily uniform) source key material  $SKM$  outputs a pseudorandom key  $PRK$ .  $\text{HKDF.Expand}(PRK, CTXinfo, L)$  on input a pseudorandom key  $PRK$  (from the Extract step) and some (potentially empty) context information  $CTXinfo$  outputs pseudorandom key material  $KM$  of length  $L$  bits. (For simplicity, we omit the third parameter  $L$  in  $\text{Expand}$  when  $L = \lambda$ , which is the case throughout TLS 1.3 except when deriving traffic keys (cf. Table 2).) Both functions are instantiated with HMAC, where directly  $\text{HKDF.Extract}(XTS, SKM) := \text{HMAC}(XTS, SKM)$  and  $\text{HKDF.Expand}$  iteratively invokes HMAC to generate pseudorandom output of the required length (see [Kra10]).

### 2.4 Dual PRF Security and the PRF-ODH Assumption

Most key derivation steps in TLS 1.3 rely on regular pseudorandom function (PRF) security for the HKDF and HMAC functions. In our analysis of the PSK handshakes, we also treat HMAC as a collision-resistant unkeyed hash function over the pair of inputs, as in Definition 2.1. For some of its applications, we however need to deploy stronger assumptions which we recap here.

The first assumption is concerned with the use of HMAC as a dual PRF (cf. [Bel06]).

**Definition 2.2** (Dual PRF security). *Let  $f: \mathcal{K} \times \mathcal{L} \rightarrow \mathcal{O}$  be a pseudorandom function with key space  $\mathcal{K}$  and label space  $\mathcal{L}$  such that  $\mathcal{K} = \mathcal{L}$ . We define the dual PRF security of  $f$  as the PRF security of  $f^{\text{swap}}(k, l) := f(l, k)$  and the according advantage function as*

$$\text{Adv}_{f, \mathcal{A}}^{\text{dual-PRF-sec}} := \text{Adv}_{f^{\text{swap}}, \mathcal{A}}^{\text{PRF-sec}}.$$

The second assumption, the so-called pseudorandom-function oracle-Diffie–Hellman (PRF-ODH) assumption, has been introduced by Jager et al. [JKSS12] in their analysis of the TLS 1.2 key exchange. It is a variant of the oracle-Diffie–Hellman assumption introduced by Abdalla et al. [ABR01] in the context of the DHIES encryption scheme. Basically, the PRF-ODH assumption states that the value  $\text{PRF}(g^{uv}, x^*)$  for a Diffie–Hellman-type key  $g^{uv}$  is indistinguishable from a random string, even when given  $g^u$  and  $g^v$  and when being able to see related values  $\text{PRF}(S^u, x)$  and/or  $\text{PRF}(T^v, x)$  for chosen values  $S, T$ , and  $x$ . The PRF-ODH assumption comes in various variants, which have been generalized and studied by Brendel et al. [BFGJ17].

For our analysis of TLS 1.3, we will deploy only the  $\text{snPRF-ODH}$  assumption providing limited oracle access to only a single related value  $\text{PRF}(S^u, x)$ , as well as its dual variant,  $\text{dual-snPRF-ODH}$ . Both have been established by Brendel et al. [BFGJ17] to hold for HMAC in the random oracle model under the strong Diffie–Hellman assumption.



**Definition 2.3** (snPRF-ODH and dual-snPRF-ODH assumptions). *Let  $\lambda \in \mathbb{N}$ ,  $\mathbb{G}$  be a cyclic group of prime order  $q$  with generator  $g$ , and  $\text{PRF}: \mathbb{G} \times \{0,1\}^* \rightarrow \{0,1\}^\lambda$  be a pseudorandom function.*

*We define the snPRF-ODH security game as follows.*

1. *The challenger samples  $b \leftarrow_{\$} \{0,1\}$ ,  $u, v \leftarrow_{\$} \mathbb{Z}_q$ , and provides  $\mathbb{G}$ ,  $g$ ,  $g^u$ , and  $g^v$  to  $\mathcal{A}$ , who responds with a challenge label  $x^*$ .*
2. *The challenger computes  $y_0^* = \text{PRF}(g^{uv}, x^*)$  and samples  $y_1^* \leftarrow_{\$} \{0,1\}^\lambda$  uniformly at random, providing  $y_b^*$  to  $\mathcal{A}$ .*
3.  *$\mathcal{A}$  may query a pair  $(S, x)$ , on which the challenger first ensures that  $S \notin \mathbb{G}$  or  $(S, x) = (g^v, x^*)$  and, if so, returns  $y \leftarrow \text{PRF}(S^u, x)$ .*
4. *Eventually,  $\mathcal{A}$  stops and outputs a guess  $b' \in \{0,1\}$ .*

*We define the snPRF-ODH advantage function as*

$$\text{Adv}_{\text{PRF}, \mathbb{G}, \mathcal{A}}^{\text{snPRF-ODH}} := 2 \cdot \Pr[b' = b] - 1.$$

*We define the dual variant of the assumption, dual-snPRF-ODH, as the snPRF-ODH assumption for a function  $\text{PRF}: \{0,1\}^* \times \mathbb{G} \rightarrow \{0,1\}^\lambda$  with swapped inputs, keyed with a group element in the second input and taking the label as first input.*

### 3 The TLS 1.3 Handshake Protocol

In this section we describe the TLS 1.3 handshake protocol modes, specifically the full one round-trip time (1-RTT) handshake, depicted on the left-hand side of Figure 1, and the combined zero round-trip time (0-RTT) and pre-shared key handshake, depicted on the right-hand side of in Figure 1. Our focus in Figure 1 and throughout the paper is on the cryptographic aspects of the TLS 1.3 handshake. As such, we omit many other components of the protocol, including most hello extensions, aspects of version and algorithm negotiation, post-handshake messages, the record layer protocol, and the alert protocol.

In TLS 1.3, the 1-RTT and PSK handshakes are divided into two distinct phases: a *key exchange phase*, where the client and the server exchange **Hello** messages to indicate support for different cryptographic options and use the selected parameters to generate key exchange material; and an *authentication phase*, where the client and the server exchange **CertificateVerify** and **Finished** messages, authenticating each other using long-term asymmetric (or symmetric) values. Figure 2 illustrates the key schedule of TLS 1.3, Table 1 lists abbreviations for messages and keys used throughout the paper, and Table 2 details some of the computations and inputs.

#### 3.1 Key-Exchange Phase

The *key exchange phase* consists of the exchange of **ClientHello** (CH) and **ServerHello** (SH) messages, during which parameters are negotiated and the core key exchange is performed, using either Diffie–Hellman key exchange or based on a pre-shared symmetric key.

**ClientHello.** The client begins by sending the **ClientHello** message, which contains  $r_c$  (a randomly-sampled 256-bit nonce value), as well as version and algorithm negotiation information.

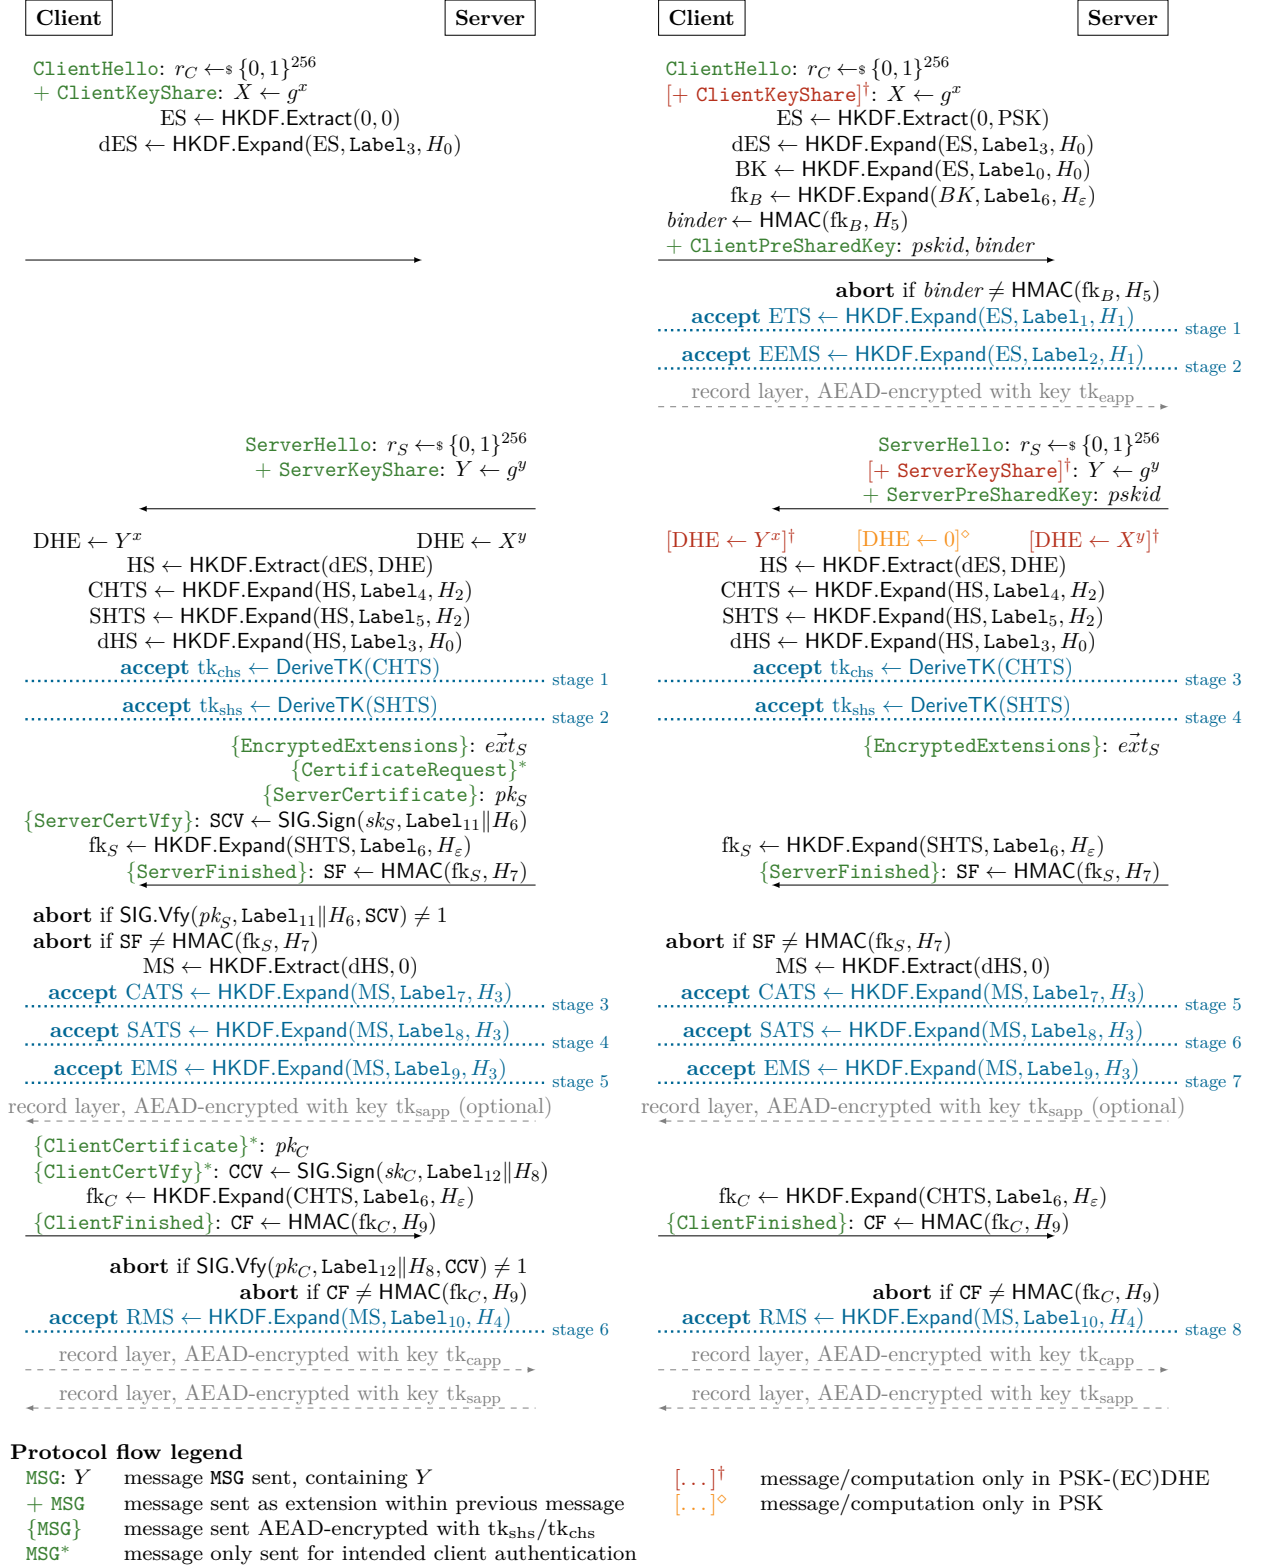


Figure 1: The TLS 1.3 full 1-RTT handshake protocol (left) and the PSK/PSK-(EC)DHE handshake protocol with optional 0-RTT (right). Shorthands are explained in Table 1; the values of context and label inputs ( $H_*$ , resp.  $\text{Label}_*$ ) and details on the calculation of traffic keys ( $tk_*$ ) can be found in Table 2.

Attached to the **ClientHello** is the **KeyShare** (CKS) extension which contains public key shares for the key exchange. Other extensions are present for further algorithm and parameter negotiation.<sup>1</sup>

If a preshared secret has been established between the client and the server (either in a previous handshake or via some out-of-band mechanism) the client may include the **PreSharedKey** (CPSK) extension, which indicates handshake modes (such as PSK or PSK-(EC)DHE) that the client supports, and a list of preshared symmetric identities that map to these PSKs.<sup>2</sup> If CPSK is included, the client computes a binder key value BK for each preshared key PSK in the list, from that a key  $fk_B$ , and a value  $binder \leftarrow \text{HMAC}(fk_B, H(\text{CH}^\dagger))$  that binds the current CH message (truncated to exclude the *binder* value itself) to each PSK, also included in the CPSK message and checked by the server. This is captured on the right-hand side of Figure 1.

Finally, if the client wishes to use the preshared secret to send zero-round-trip time (0-RTT) data, the client can indicate this by sending a **EarlyDataIndication** extension. This will indicate to the server that the client will use the first preshared secret indicated in the CPSK list to derive an early traffic secret (ETS), and early exporter master secret (EEMS), and begin sending encrypted data to the server without first requiring the client to receive **ServerHello** response.

**ServerHello.** The next message in the key-exchange phase is the **ServerHello** (SH) message. As in CH, the server will randomly sample a 256-bit nonce value  $r_s$ . The server picks among the various algorithms and parameters offered by the client and responds with its selections. If CPSK was sent, the server decides whether to accept a PSK-based handshake. If so, then the preshared key identifier *pskid* associated with the selected PSK is sent in the **PreSharedKey** (SPSK) extension. If the server has chosen PSK-(EC)DHE mode (or has rejected the use of PSKs), the server will generate its own (EC)DHE key share  $Y \leftarrow g^y$ , sending  $Y$  in the **KeyShare** (SKS) extension attached to SH.

At this point, the server has enough information to compute the client handshake traffic secret (CHTS) and server handshake traffic secret (SHTS) values, and uses these to derive client and server handshake traffic keys ( $tk_{chs}$  and  $tk_{shs}$ , respectively). The first part of Figure 2 shows the key schedule for deriving these keys. Note that we consider  $tk_{chs}$  and  $tk_{shs}$  being derived at the same point in time (namely when the handshake secret HS becomes available), although  $tk_{chs}$  is in principle only needed a bit later.

The server now begins to encrypt all handshake messages under  $tk_{shs}$ , and any extensions that are not required to establish the server handshake traffic key are sent (and encrypted) in the **EncryptedExtensions** (EE) messages.

### 3.2 Authentication Phase

The *authentication phase* now begins. All handshake messages in this phase are encrypted under  $tk_{shs}$  or  $tk_{chs}$ . In the full 1-RTT handshake, authentication is based on public key certificates; see the left-hand side of Figure 1. In pre-shared key handshakes (both PSK and PSK-(EC)DHE), the server and client will authenticate each other by relying on a message authentication code applied to the transcript; see the right-hand side of Figure 1.

<sup>1</sup>Note that our analysis in Sections 5 and 6 does not consider the negotiation of cryptographic values (such as preshared keys or (EC)DHE groups) or handshake modes, but instead our analysis considers each handshake mode and ciphersuite combination in isolation. This can be seen in Figure 1, e.g., the CKS message contains only a single (EC)DHE key share.

<sup>2</sup>As for the **KeyShare** extension, we do not consider negotiation here and only capture the single PSK entry that client and server agree upon.

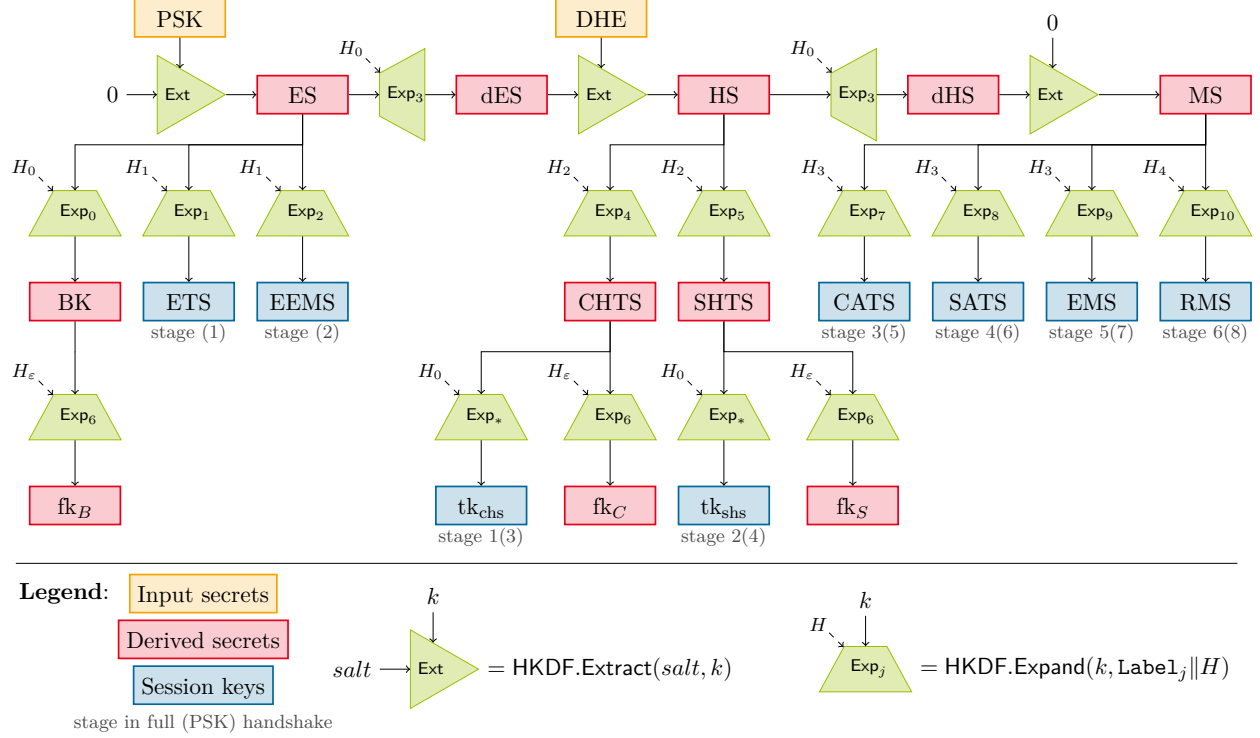


Figure 2: The TLS 1.3 key schedule. The values of context and label inputs ( $H_*$ , resp.  $\text{Label}_*$ ) and details on the calculation of traffic keys ( $\text{Exp}_*$ ) can be found in Table 2.

**Authentication in full 1-RTT handshake.** The server can request public-key-based client authentication by sending a **CertificateRequest** (CR) message. The server will authenticate to the client by using the server’s long-term public keys. Here, the server begins by sending its certificate (carrying its public key) in the **ServerCertificate** (SCRT) message. The server then computes **ServerCertificateVerify** authentication value by signing the session hash (which is a continuously updating hash of all messages up to this point in the protocol), then sends it to the client as the **ServerCertificateVerify** message.

**Server key confirmation and key derivation.** In all handshake modes, the final message that the server sends to the client is the **ServerFinished** (SF) message. The server first derives a server finished key  $\text{fk}_S$  from SHTS and then computes a MAC tag SF over the session hash. This value is also encrypted under  $\text{tk}_{\text{shs}}$ , sending the output ciphertext to the client. At this point, the server is able to compute the client application traffic secret (CATS), the server application traffic secret (SATS), and the exporter master secret (EMS). Figure 2 shows the key schedule for deriving these keys and all other keys in the authentication phase. Now that the server has computed the server application traffic key  $\text{tk}_{\text{sapp}}$ , it can begin sending encrypted application data to the client without waiting for the final flight of messages from the client, thus achieving a 0.5-RTT handshake.

**Client verification, authentication, key confirmation, and key derivation.** The client, upon receiving these messages, checks that the signature SCV (if in full 1-RTT mode) and the MAC SF verify correctly. If the server has requested client authentication, the client will begin by sending its digital certificate (carrying its public-key) in the **ClientCertificate** (CCRT) message,

Message		Derived key or value	
CH	ClientHello	BK	Binder Key
CKS	ClientKeyShare	CHTS/SHTS	Client/Server Handshake Traffic Secret
CPSK	ClientPreSharedKey	CATS/SATS	Client/Server Application Traffic Secret
SH	ServerHello	dES/dHS	Derived Early/Handshake Secret
SKS	ServerKeyShare	ES/HS/MS	Early/Handshake/Master Secret
SPSK	ServerPreSharedKey	ETS	Early Traffic Secret
EE	EncryptedExtensions	EEMS/EMS	(Early) Exporter Master Secret
CR	CertificateRequest	$fk_B/fk_C/fk_S$	Binder/Client/Server Finished Key
SCRT	ServerCertificate	RMS	Resumption Master Secret
SCV	ServerCertificateVerify	$tk_{eapp}$	Early Application Traffic Key
SF	ServerFinished	$tk_{chs}/tk_{shs}$	Client/Server Handshake Traffic Key
CCRT	ClientCertificate	$tk_{capp}/tk_{sapp}$	Client/Server Application Traffic Key
CCV	ClientCertificateVerify		
CF	ClientFinished		

Table 1: Shorthands for TLS 1.3 messages (in protocol order) and derived keys/values (alphabetical).

after which the client will compute its own certificate verify value **CCV** by signing the session hash, then send it to the server as the **CCV** message. The client finally derives the client finished key  $fk_C$  from CHTS and uses  $fk_C$  to compute a MAC tag **CF** over the session hash.

**Server verification.** The server will verify the final MAC (**SF**) and optional signature (**SCV**) messages of the client.

**Handshake completion.** At this point both parties can compute the resumption master secret (RMS) value that can be used as a pre-shared key for session resumption in the future. Both parties can now derive the client application traffic key ( $tk_{capp}$ ), and use the record layer for encrypted communication of application data with the resulting keys.

### 3.3 NewSessionTicket

The **NewSessionTicket** message is a post-handshake message in TLS 1.3 which refers to values from the handshake protocol. The **NewSessionTicket** message can be sent by a server to the client (encrypted under a server application traffic key  $tk_{sapp}$ ) to allow the client to compute values associated with resumption handshakes, including the PSK used in resumption as well as an identifier to indicate to the server which pre-shared key is being used. The **NewSessionTicket** message contains two fields that are interesting for this purpose:

- **ticket\_nonce**, which is used by the client as the salt value to derive the pre-shared key to be used in future handshake for resumption:  $PSK \leftarrow \text{HKDF.Expand}(\text{RMS}, \text{"resumption"}, \text{ticket\_nonce})$ .
- **ticket**, which is an opaque label used to publicly refer to the associated preshared key in future **PreSharedKey** messages. In our notation used in Figure 1, the preshared key identifier  $pskid = \text{ticket}$ .

In our analysis, we do not capture this **NewSessionTicket** message, nor the derivation of PSK from RMS, and instead assume that the mapping between PSK and  $pskid$  is established in some out-of-band way. In particular, we do not capture transmission of **NewSessionTicket** under a server application traffic key  $tk_{sapp}$ , as it would impact how we consider the usage of SATS. In our analysis, we currently consider SATS an “external key” used in an arbitrary symmetric-key

Secret	Context Input	Label Input
BK	$H_0 = H("")$	Label <sub>0</sub> = "ext binder" / "res binder"
fk <sub>B</sub>	$H_\varepsilon = ""$	Label <sub>6</sub> = "finished"
ETS	$H_1 = H(\text{ClientHello})$	Label <sub>1</sub> = "c e traffic"
EEMS	$H_1 = H(\text{ClientHello})$	Label <sub>2</sub> = "e exp master"
dES	$H_0 = H("")$	Label <sub>3</sub> = "derived"
CHTS	$H_2 = H(\text{ClientHello} \parallel \text{ServerHello})$	Label <sub>4</sub> = "c hs traffic"
SHTS	$H_2 = H(\text{ClientHello} \parallel \text{ServerHello})$	Label <sub>5</sub> = "s hs traffic"
fk <sub>S</sub>	$H_\varepsilon = ""$	Label <sub>6</sub> = "finished"
dHS	$H_0 = H("")$	Label <sub>3</sub> = "derived"
CATS	$H_3 = H(\text{ClientHello} \parallel \dots \parallel \text{ServerFinished})$	Label <sub>7</sub> = "c ap traffic"
SATS	$H_3 = H(\text{ClientHello} \parallel \dots \parallel \text{ServerFinished})$	Label <sub>8</sub> = "s ap traffic"
EMS	$H_3 = H(\text{ClientHello} \parallel \dots \parallel \text{ServerFinished})$	Label <sub>9</sub> = "exp master"
fk <sub>C</sub>	$H_\varepsilon = ""$	Label <sub>6</sub> = "finished"
RMS	$H_4 = H(\text{ClientHello} \parallel \dots \parallel \text{ClientFinished})$	Label <sub>10</sub> = "res master"
Auth. Value	Context Input	Context String (for signatures only)
<i>binder</i>	$H_5 = H(\text{ClientHello}^\dagger)$	
SCV	$H_6 = H(\text{ClientHello} \parallel \dots \parallel \text{ServerCert})$	Label <sub>11</sub> = "TLS 1.3, server CertificateVerify"
SF	$H_7 = H(\text{ClientHello} \parallel \dots \parallel \text{ServerCertVfy})$	
CCV	$H_8 = H(\text{ClientHello} \parallel \dots \parallel \text{ClientCert})$	Label <sub>12</sub> = "TLS 1.3, client CertificateVerify"
CF	$H_9 = H(\text{ClientHello} \parallel \dots \parallel \text{ClientCertVfy}^*)$	
Traffic Key Calculation		
$\text{tk}_{\text{eapp}}/\text{tk}_{\text{chs}}/\text{tk}_{\text{shs}}/\text{tk}_{\text{capp}}/\text{tk}_{\text{sapp}} = (\text{key}, \text{iv}) = \text{DeriveTK}(\text{ETS}/\text{CHTS}/\text{SHTS}/\text{CATS}/\text{SATS})$ where $\text{DeriveTK}(\text{Secret}) = (\text{HKDF.Expand}(\text{Secret}, \text{"key"}, H(""), L_k), \text{HKDF.Expand}(\text{Secret}, \text{"iv"}, H(""), L_{iv}))$ with $L_k/L_{iv}$ indicating the key/iv length of the negotiated AEAD scheme		

Table 2: Secret, context, and label inputs to the HKDF.Expand resp. authentication functions as well as traffic key calculation in the TLS 1.3 handshake (Figure 1) and key schedule (Figure 2). The actual label input to HKDF.Expand is the concatenation of the hash length (in bytes), the string "tls13 ", Label, and the given context value. HKDF.Expand is then called on the corresponding secret, this augmented label, and the desired output length. ClientCertVfy\* is only included in case of client authentication. ClientHello<sup>†</sup> indicates a truncated version of ClientHello which excludes the binder value itself. Signatures in SCV and CCV are computed over the concatenation of a constant (0x20 repeated 64 times), the label as context information, a separating 0 byte, and the context value.

protocol. To capture the transmission of NewSessionTicket, we would need to capture the use of SATS in deriving tk<sub>sapp</sub> and then establishing PSK. We choose to simplify the analysis by omitting this mechanism, and leave this as future work.

## 4 Multi-Stage Key Exchange Security Model

In order to capture the security of all variants of the TLS 1.3 handshake within a single comprehensive security model, we adopt the multi-stage key exchange model in the version by Günther [Gün18] which combines the original model by Fischlin and Günther [FG14] with follow-up extensions [DFGS15, DFGS16, FG17]. We refer to Günther [Gün18] for an extensive discussion of the model, but recap its core concepts and definitions as well as adaptations for our analysis in the following.

The model follows the classical paradigm for key exchange models of Bellare and Rogaway [BR94] in the formalism of Brzuska et al. [BFWW11, Brz13]. This paradigm captures a strong adversary that controls the network and is able to both passively eavesdrop and to actively modify the com-



munication across multiple sessions of the key exchange protocol (spawning them via a **NewSession** oracle and directing communication via a **Send** oracle). The adversary is further allowed to expose the long-term secrets of interacting honest parties (via a **Corrupt** oracle) as well as the session keys in some protocol runs (through a **Reveal** oracle). Basic security then demands that such adversary cannot distinguish the real established session key in some uncompromised (“fresh”) session from a random one (through a **Test** oracle).

The multi-stage key exchange model now extends the basic key exchange setting by capturing protocols that derive a series of session keys in multiple *stages*. Each stage is associated with particular security properties, steering admissibility of certain adversarial actions for that stage and under which conditions the key of this stage is considered fresh. These security properties model the following aspects:

**Authentication.** Our model distinguishes between *unauthenticated* stages, *unilaterally authenticated* stages where only the responder (the server in TLS 1.3) authenticates, and *mutually authenticated* stages where both peers authenticate. We treat the authentication of each stage *individually* and consider *concurrent* executions of different authentication modes of the same protocol. The identities of communication partners may be learned only during the execution of the protocol (e.g., through exchanged certificates), which we implement through *post-specified peers* following Canetti and Krawczyk [CK02]. Our model demands a strong notion of security for sessions with unauthenticated peers, namely that such sessions achieve key secrecy when receiving their messages from an honest session (identified via a *contributive identifier*), independent of whether that honest peer session later becomes partnered.

Moreover, the authentication level of some stage may be raised with acceptance of a later stage, e.g. from unauthenticated to unilaterally or even mutually authenticated. This may happen for instance if a party later signs previously transmitted data, as in case of TLS 1.3. We capture this by allowing a protocol to specify the authentication level for each acceptance stage, as well as at which later stage(s) that level increases.

Note that we capture authentication *implicitly* through key secrecy (i.e., keys are only known to the intended peer session) but do not prove explicit authentication (i.e., the existence of a partnered session). The SIGMA design [Kra03], on which the main TLS 1.3 handshake is based, ensures explicit authentication. de Saint Guilhem et al. [dFW19] give a generic argument that explicit authentication follows from implicit key secrecy (which is shown for TLS 1.3 in this article) and key confirmation [FGSW16].

**Forward secrecy.** We capture the usual notion of forward secrecy, which ensures that accepted session keys remain secure after a long-term secret compromise. In a multi-stage key exchange protocol, forward secrecy may however be reached only from a certain stage on (e.g., due to mixing-in forward-secret key material). The model hence treats *stage- $j$  forward secrecy*, indicating that keys from stage  $j$  on are forward secret.

**Key usage.** Some stage keys might be used internally in the key exchange protocol, e.g., in the case of TLS 1.3 the handshake key is used to encrypt part of the key exchange communication. We distinguish the usage of keys as *internal* when used within the key exchange, and *external* when exclusively used outside of the key exchange (e.g., to encrypt application data). In the former case, our model ensures that tested real-or-random keys are accordingly used in subsequent key exchange steps, and pauses the protocol execution to enable testing of those keys. We note that the declaration of whether a key is internal or external is a parameter to the model, and becomes a part of the protocol description and its security guarantees.

**Public or pre-shared keys.** Our multi-stage model comes in two flavors that capture both the regular, public-key case (abbreviated as pMSKE) of long-term keys being public/secret key pairs (as in the TLS 1.3 full handshake) as well as the pre-shared-secret case (abbreviated sMSKE) case where pre-shared symmetric keys act as long-term secrets (as in the TLS 1.3 resumption handshake).

**Replayability.** For 0-RTT key establishment, key exchange protocols (including TLS 1.3) regularly give up strong replay protection guarantees, in the sense that client (initiator) messages can be replayed to several server (responder) sessions. We capture this in our model by distinguishing between *replayable* (0-RTT) and regular *non-replayable* stages, taking potential replays into account for the former while still requiring key secrecy. Determining the replay type of a stage is again a parameter to the model and must be specified as part of the protocol description resp. the security claim.

We note that former variants of multi-stage key exchange models including [Gün18] further differentiated whether the compromise of some stage’s key affects the security of other stages’ keys under the notion of *key (in)dependence*. Here, we always demand such compromise never affects other stages’ keys as the desirable goal, i.e., we postulate key independence and reduce the model’s complexity by incorporating this property straight into the model. As we will see, TLS 1.3 always achieves this property due to clean key separation in the key scheduling, and already did so in earlier draft versions [DFGS15, DFGS16, FG17].

**Secret compromise paradigm.** We follow the paradigm of the Bellare–Rogaway model [BR94], focusing on the leakage of long-term secret inputs and session key outputs of the key exchange, but not on internal values within the execution of a session. This contrasts to some extent with the model by Canetti and Krawczyk [CK01] resp. LaMacchia et al. [LLM07] which include a “session state reveal” resp. “ephemeral secret reveal” query that allows accessing internal variables of the session execution.

In the context of TLS 1.3, this means we consider the leakage of:

- *Long-term keys* (such as the signing keys of the server or client, but also their pre-shared keys), since long-term values have the potential to be compromised, and this is necessary to model forward secrecy; it is captured in our model by the **Corrupt** query.
- *Session keys* (such as the various traffic encryption keys or the derived resumption or exporter secrets), since these are outputs of the key exchange and are used beyond this protocol for encryption, later resumption, or exporting of keying material; this is modeled by the **Reveal** query.

We do not permit the leakage of:

- *Ephemeral secrets / randomness* (such as the randomness in a signature algorithm or ephemeral Diffie–Hellman exponents); this is disallowed since TLS 1.3 is not designed to be secure if these values are compromised.
- *Internal values / session state* (e.g., internally computed master secrets or MAC keys); this is disallowed since TLS 1.3 is not designed to be secure if these values are compromised.

**Comparison with previous multi-stage key exchange models.** Compared to the original MSKE model of Fischlin and Günther [FG14], the most notable changes in our model are the addition which models upgradeable authentication and accommodating both public and pre-shared symmetric keys for authentication. We also do not track whether keys are independent or not, as all keys established in TLS 1.3 satisfy key independence (unlike in the analysis of QUIC in [FG14]). Key usage (internal versus external) and replayability were introduced to MSKE by [FG17].

## 4.1 Syntax

In our model, we explicitly separate some *protocol-specific* properties (as, e.g., various authentication flavours) from *session-specific* properties (as, e.g., the state of a running session). We represent protocol-specific properties as a vector  $(M, \text{AUTH}, \text{FS}, \text{USE}, \text{REPLAY})$  that captures the following:

- $M \in \mathbb{N}$ : the number of stages (i.e., the number of keys derived).<sup>3</sup>
- $\text{AUTH} \subseteq \{((u_1, m_1), \dots, (u_M, m_M)) \mid u_j, m_j \in \{1, \dots, M, \infty\}\}$ : a set of vectors of pairs, each vector encoding a supported scheme for authentication and authentication upgrades, for each stage. For example, the  $i$ -th entry  $(u_i, m_i)$  in a vector says that the session key in stage  $i$  initially has the default *unauthenticated* level, i.e., provides no authentication for either communication partner, then at stage  $u_i$  becomes *unilaterally authenticated* and thus authenticates only the responder (server), and becomes *mutually authenticated* to authenticate both communication partners at stage  $m_j$ . Note that we allow for example  $u_i = i$  (or even  $u_i = m_i = i$ ) such that the session key is immediately unilaterally (resp. mutually) authenticated when derived.

Entries in each pair must be non-decreasing, and  $u_i = \infty$  or  $m_i = \infty$  denotes that unilateral, resp. mutual, authentication is never reached for stage  $i$ .

- $\text{FS} \in \{1, \dots, M, \infty\}$ : the stage  $j = \text{FS}$  from which on keys are forward secret (or  $\infty$  in case of no forward secrecy).<sup>4</sup>
- $\text{USE} \in \{\text{internal}, \text{external}\}^M$ : the usage indicator for each stage, where  $\text{USE}_i$  indicates the usage of the stage- $i$  key. Here, an internal key is used within the key exchange protocol (but possibly also externally), whereas an external key must not be used within the protocol, making the latter potentially amenable to generic composition (cf. Section 7.3). As shorthand notation, we, e.g., write  $\text{USE} = (\text{internal} : \{1, 4\}, \text{external} : \{2, 3, 5\})$  to indicate that usage of keys in stage 1 and 4 is internal, and external for the other stages.
- $\text{REPLAY} \in \{\text{replayable}, \text{nonreplayable}\}^M$ : the replayability indicator for each stage, where  $\text{REPLAY}_i$  indicates whether the  $i$ -th stage is replayable in the sense that an adversary can easily force identical communication and thus identical session identifiers and keys in this stage (e.g., by re-sending the same data in 0-RTT stages). Note that the adversary, however, should still not be able to distinguish such a replayed key from a random one. We remark that, from a security viewpoint, the usage of replayable stages should ideally be limited, although such stages usually come with an efficiency benefit. We use the same shorthand

<sup>3</sup>We fix a maximum stage  $M$  only for ease of notation. Note that  $M$  can be arbitrarily large in order to cover protocols where the number of stages is not bounded a-priori. Also note that stages and session key derivations may be “back to back,” without further protocol interactions between parties.

<sup>4</sup>A more general multi-stage key exchange model could have a vector tracking specifically which subset of stage keys have forward secrecy. We do not need such generality since forward secrecy is monotonic in TLS 1.3.

notation as for USE; e.g.,  $\text{REPLAY} = (\text{nonreplayable} : \{1, 2, 3\})$  indicates that all three stages are non-replayable.

We denote by  $\mathcal{U}$  the set of *identities* (or *users*) used to model the participants in the system, each identified by some  $U \in \mathcal{U}$ . *Sessions* of a protocol are uniquely identified (on the administrative level of the model) using a *label*  $\text{label} \in \text{LABELS} = \mathcal{U} \times \mathcal{U} \times \mathbb{N}$ , where  $\text{label} = (U, V, n)$  indicates the  $n$ -th local session of identity  $U$  (the session *owner*) with  $V$  as the intended communication *partner*.

In the public-key variant of the model (pMSKE), each identity  $U$  is associated with a certified *long-term* public key  $\text{pk}_U$  and secret key  $\text{sk}_U$ . In the pre-shared secret setting (sMSKE), a session instead holds an identifier  $\text{pssid} \in \{0, 1\}^*$  for the pre-shared secret  $\text{pss} \in \mathcal{P}$  (from some pre-shared secret space  $\mathcal{P}$ ) used. The challenger maintains maps  $\text{pss}_{U,V} : \{0, 1\}^* \rightarrow \mathcal{P}$  mapping an identifier to the corresponding secret shared by parties  $U$  and  $V$ , where  $U$  uses that secret (only) in the initiator role and  $V$  (only) in the responder role<sup>5</sup>, and for any user  $U$ , a pre-shared secret identifier  $\text{pssid}$  uniquely identifies the peer identity  $V$  it is shared with.

For each session, a tuple with the following information is maintained as an entry in the *session list*  $\text{List}_S$ , where values in square brackets  $[]$  indicate the default initial value. Some variables have values for each stage  $i \in \{1, \dots, M\}$ .

- $\text{label} \in \text{LABELS}$ : the unique (administrative) session label
- $\text{id} \in \mathcal{U}$ : the identity of the session owner
- $\text{pid} \in \mathcal{U} \cup \{*\}$ : the identity of the intended communication partner, where the distinct wildcard symbol ‘ $*$ ’ stands for “currently unknown identity” but can be later set to a specific identity in  $\mathcal{U}$  once by the protocol
- $\text{role} \in \{\text{initiator}, \text{responder}\}$ : the session owner’s role in this session
- $\text{auth} \in \text{AUTH}$ : the intended authentication type vector from the set of supported authentication properties  $\text{AUTH}$ , where  $\text{auth}_i$  indicates the authentication level pair for stage  $i$ , and  $\text{auth}_{i,j}$  its  $j$ -th entry
- $\text{pssid} \in \{0, 1\}^* \cup \{\perp\}$ : In the pre-shared secret (sMSKE) variant the identifier for the pre-shared secret (i.e.,  $\text{pss}_{\text{id}, \text{pid}}$  if  $\text{role} = \text{initiator}$ , else  $\text{pss}_{\text{pid}, \text{id}}$ ) to be used in the session; can be initialized with  $\perp$  if  $\text{pid} = *$  is unknown and then must be set (once) when  $\text{pid}$  is set
- $\text{st}_{\text{exec}} \in (\text{RUNNING} \cup \text{ACCEPTED} \cup \text{REJECTED})$ : the state of execution  $[\text{running}_0]$ , where  $\text{RUNNING} = \{\text{running}_i \mid i \in \mathbb{N} \cup \{0\}\}$ ,  $\text{ACCEPTED} = \{\text{accepted}_i \mid i \in \mathbb{N}\}$ ,  $\text{REJECTED} = \{\text{rejected}_i \mid i \in \mathbb{N}\}$ ; set to  $\text{accepted}_i$  in the moment a session accepts the  $i$ -th key, to  $\text{rejected}_i$  when the session rejects that key (a session may continue after rejecting in a stage<sup>6</sup>), and to  $\text{running}_i$  when a session continues after accepting the  $i$ -th key
- $\text{stage} \in \{0, \dots, M\}$ : the current stage  $[0]$ , where  $\text{stage}$  is incremented to  $i$  when  $\text{st}_{\text{exec}}$  reaches  $\text{accepted}_i$  resp.  $\text{rejected}_i$
- $\text{sid} \in (\{0, 1\}^* \cup \{\perp\})^M$ :  $\text{sid}_i [\perp]$  indicates the session identifier in stage  $i$ , set once (and only) upon acceptance in that stage

<sup>5</sup>Requiring a fixed role in which a pre-shared key can be used by either peer avoids the Selfie attack [DG21b, AASS19].

<sup>6</sup>This models, e.g., servers rejecting 0-RTT data from a client, but continuing with the remaining handshake.

- $\text{cid} \in (\{0, 1\}^* \cup \{\perp\})^M$ :  $\text{cid}_i [\perp]$  indicates the contributive identifier in stage  $i$ , may be set several times until acceptance in that stage
- $\text{key} \in (\{0, 1\}^* \cup \{\perp\})^M$ :  $\text{key}_i [\perp]$  indicates the established session key in stage  $i$ , set once upon acceptance in that stage
- $\text{st}_{\text{key}} \in \{\text{fresh}, \text{revealed}\}^M$ :  $\text{st}_{\text{key}, i} [\text{fresh}]$  indicates the state of the session key in stage  $i$
- $\text{tested} \in \{\text{true}, \text{false}\}^M$ : test indicator  $\text{tested}_i [\text{false}]$ , where **true** means that  $\text{key}_i$  has been tested
- $\text{corrupted} \in \{0, \dots, M, \infty\}$ : corruption indicator  $[\infty]$  holding the stage the session was in when a **Corrupt** was issued to its owner or intended partner, including the value 0 if the corruption had taken place before the session started, and  $\infty$  if none of the parties is corrupted

By convention, adding a not fully specified tuple  $(\text{label}, \text{id}, \text{pid}, \text{role}, \text{auth})$  resp.  $(\text{label}, \text{id}, \text{pid}, \text{role}, \text{auth}, \text{pssid})$  to  $\text{List}_S$  sets all other entries to their default value. As shorthands, for some tuple with (unique) label  $\text{label}$  in  $\text{List}_S$  we furthermore write  $\text{label}.X$  for that tuple's element  $X$  and  $\text{label}.(X, Y, Z)$  for the vector  $(X, Y, Z)$  of that tuple's elements  $X$ ,  $Y$ , and  $Z$ .

We define two distinct sessions  $\text{label}$  and  $\text{label}'$  to be *partnered* in stage  $i$  if both sessions hold the same session identifier in that stage, i.e.,  $\text{label}.\text{sid}_i = \text{label}'.\text{sid}_i \neq \perp$ , and require for correctness that two sessions having a non-tampered joint execution are partnered in all stages upon acceptance.

Our security model treats corruption of long-term secrets (secret keys for pMSKE, pre-shared secrets for sMSKE). While the affects of such compromises on sessions may differ in each setting, we broadly consider the derived keys of some session to be revealed if, in the public-key setting (pMSKE), the owner or peer secret key is compromised, or in the pre-shared secret setting (sMSKE), if the pre-shared secret used for that session is compromised. Forward secrecy comes into play when determining if keys derived prior to the long-term secret corruption are affected, too. In more precise notation, we say a session  $\text{label}$  is *corrupted* if

- for pMSKE, the session's owner  $\text{label}.\text{id}$  or intended communication partner  $\text{label}.\text{pid}$  is corrupted (i.e.,  $\{\text{label}.\text{id}, \text{label}.\text{pid}\} \cap \mathcal{C} \neq \emptyset$ ), resp.
- for sMSKE, the used pre-shared secret is corrupted (i.e.,  $(\text{label}.\text{id}, \text{label}.\text{pid}, \text{label}.\text{pssid}) \in \mathcal{C}$ , the set of corrupted users) if  $\text{label}.\text{role} = \text{initiator}$ , resp.  $(\text{label}.\text{pid}, \text{label}.\text{id}, \text{label}.\text{pssid}) \in \mathcal{C}$  if  $\text{label}.\text{role} = \text{responder}$ .

**Upgradable authentication.** We capture that the authentication level of some stage may increase, possibly twice, with acceptance of a later stage through a per-stage vector in the authentication level matrix. When capturing security, our model however needs to carefully consider the interaction of authentication and corruptions (somewhat similar to what one might be used to for forward secrecy). More precisely, the authentication guarantee of some stage  $i$  *after its acceptance* can only step up (in some later stage  $j > i$ ) if the involved parties are not corrupted by the time stage  $j$  accepts. Otherwise, the adversary may have impersonated the party up to the unauthenticated stage  $i$  and now post-authenticates as the party after corruption in stage  $j$ . This would effectively mean that the adversary has been in full control of the session and may thus know the session key of stage  $i$ .

We capture the upgrade by defining the *rectified authentication level*  $\text{rect\_auth}_i$  of some stage  $i$  in a session with intended authentication vector  $\text{auth}$ , consisting of pairs  $(\text{auth}_{i,1}, \text{auth}_{i,2})$  describing

the stage in which the  $i$ -th session key gets unilaterally and mutually authenticated, with corruption indicator `corrupted`, and with current execution stage `stage` as follows:

$$\text{rect\_auth}_i := \begin{cases} \text{mutual} & \text{if stage} \geq \text{auth}_{i,2} \text{ and corrupted} \geq \text{auth}_{i,2} \\ \text{unilateral} & \text{if stage} \geq \text{auth}_{i,1} \text{ and corrupted} \geq \text{auth}_{i,1} \\ \text{unauth} & \text{otherwise} \end{cases}$$

This encodes that authentication level of stage  $i$  is upgraded (to unilateral or mutual) when reaching stage `authi,1`, resp. `authi,2`, only if no corruption affected this session prior to these stages (`authi,1`, resp. `authi,2`).

## 4.2 Adversary Model

We consider a probabilistic polynomial-time (PPT) adversary  $\mathcal{A}$  which controls the communication between all parties, enabling interception, injection, and dropping of messages. Our adversary model further reflects the advanced security aspects in multi-stage key exchange as outlined above. We conveniently capture admissibility of adversarial interactions and conditions where the adversary trivially loses (such as both revealing and testing the session key in partnered sessions) via a flag `lost` (initialized to `false`).

The adversary interacts with the protocol via the following queries.

- **NewSecret**( $U, V, \text{pssid}$ ): This query is only available in the pre-shared secret (sMSKE) variant. Generates a fresh secret with identifier `pssid` shared between parties  $U$  and  $V$ , to be used by  $U$  in the initiator role and by  $V$  in the responder role. If `pssU,V(pssid)` is already set, return  $\perp$  to ensure uniqueness of `pssid` identifiers between two parties in these roles. Otherwise, sample `pss`  $\leftarrow^{\$} \mathcal{P}$  uniformly at random from the protocol's pre-shared secret space  $\mathcal{P}$  and define `pssU,V(pssid) := pss`.

- **NewSession**( $U, V, \text{role}, \text{auth}[, \text{pssid}]$ ): Creates a new session with a (unique) new label `label` for owner identity `id = U` with role `role`, having `pid = V` as intended partner (potentially unspecified, indicated by  $V = *$ ) and aiming at authentication type `auth`.

In the pre-shared secret (sMSKE) variant, the additional parameter `pssid` identifies the pre-shared secret to be used, namely `pssU,V(pssid)` if `role = initiator`, resp. `pssV,U(pssid)` if `role = responder`. The identifier might be unspecified at this point (indicated by `pssid =  $\perp$` ) and may then be set later by the protocol once.

Add `(label, U, V, role, auth)`, resp. `(label, U, V, role, auth, pssid)`, to `ListS`. If `label` is corrupted, set `label.corrupted`  $\leftarrow 0$ . This encodes the information that the session is corrupt right from the beginning. Return `label`.

- **Send**(`label, m`): Sends a message  $m$  to the session with label `label`.

If there is no tuple with label `label` in `ListS`, return  $\perp$ . Otherwise, run the protocol on behalf of  $U$  on message  $m$  and return the response and the updated state of execution `label.stexec`. As a special case, if `label.role = initiator` and  $m = \text{init}$ , the protocol is initiated (without any input message).

If, during the protocol execution, the state of execution changes to `acceptedi`, the protocol execution is immediately suspended and `acceptedi` is returned as result to the adversary. The adversary can later trigger the resumption of the protocol execution by issuing a special



Send(label, continue) query. For such a query, the protocol continues as specified, with the party creating the next protocol message and handing it over to the adversary together with the resulting state of execution  $\text{st}_{\text{exec}}$ . We note that this is necessary to allow the adversary to test an internal key, before it may be used immediately in the response and thus cannot be tested anymore to prevent trivial distinguishing attacks. It furthermore allows the adversary to corrupt long-term keys in a fine-grained manner after any acceptance of a key.

If the state of execution changes to  $\text{label.st}_{\text{exec}} = \text{accepted}_i$  for some  $i$  and there is a partnered session  $\text{label}' \neq \text{label}$  in  $\text{List}_S$  (i.e.,  $\text{label.sid}_i = \text{label'.sid}_i$ ) with  $\text{label'.tested}_i = \text{true}$ , then set  $\text{label.tested}_i \leftarrow \text{true}$  and (only if  $\text{USE}_i = \text{internal}$ )  $\text{label.key}_i \leftarrow \text{label'.key}_i$ . This ensures that, if the partnered session has been tested before, subsequent Test queries for the session are answered accordingly and, in case it is used internally, this session's key  $\text{key}_i$  is set consistently.<sup>7</sup>

If the state of execution changes to  $\text{label.st}_{\text{exec}} = \text{accepted}_i$  for some  $i$  and the session label is corrupted, then set  $\text{label.st}_{\text{key},i} \leftarrow \text{revealed}$ .

- **Reveal(label, i):** Reveals the session key  $\text{label.key}_i$  of stage  $i$  in the session with label label. If there is no session with label label in  $\text{List}_S$  or  $\text{label.stage} < i$ , then return  $\perp$ . Otherwise, set  $\text{label.st}_{\text{key},i}$  to revealed and provide the adversary with  $\text{label.key}_i$ .

- **Corrupt( $U$ )** or **Corrupt( $U, V, \text{pssid}$ )**: The first query is only used in the public-key (pMSKE) variant, the second query only in the pre-shared secret (sMSKE) variant. Provide the adversary with the corresponding long-term secret, i.e.,  $\text{sk}_U$  (pMSKE), resp.  $\text{pss}_{U,V}(\text{pssid})$  (sMSKE). Add to the set of corrupted entities  $\mathcal{C}$  the user  $U$  (for pMSKE), resp. (for sMSKE) the global pre-shared secret identifier  $(U, V, \text{pssid})$ .

Record the time of corruption in each session label with  $\text{label.id} = U$  or  $\text{label.pid} = U$  (pMSKE), resp. with  $\text{label}(\text{role}, \text{id}, \text{pid}, \text{pssid}) \in \{(\text{initiator}, U, V, \text{pssid}), (\text{responder}, V, U, \text{pssid})\}$  (sMSKE), by setting  $\text{label.corrupted} \leftarrow \text{label.stage}$  (unless  $\text{label.corrupted} \neq \infty$  already, in which case corruption had taken place earlier such that we leave the value unchanged).

In the non-forward-secret case, for each such session label and for all  $i \in \{1, \dots, M\}$ , set  $\text{label.st}_{\text{key},i}$  to revealed. I.e., all (previous and future) session keys are considered to be disclosed.

In the case of stage- $j$  forward secrecy,  $\text{st}_{\text{key},i}$  of each such session label is instead set to revealed only if  $i < j$  or if  $i > \text{stage}$ . This means that session keys before the  $j$ -th stage (where forward secrecy kicks in) as well as keys that have not yet been established are potentially disclosed.

- **Test(label, i):** Tests the session key of stage  $i$  in the session with label label. In the security game this oracle is given a uniformly random test bit  $b_{\text{test}}$  as state which is fixed throughout the game.

If there is no session with label label in  $\text{List}_S$  or if  $\text{label.st}_{\text{exec}} \neq \text{accepted}_i$  or  $\text{label.tested}_i = \text{true}$ , return  $\perp$ . If stage  $i$  is internal (i.e.,  $\text{USE}_i = \text{internal}$ ) and there is a partnered session  $\text{label}'$  in  $\text{List}_S$  (i.e.,  $\text{label.sid}_i = \text{label'.sid}_i$ ) with  $\text{label'.st}_{\text{exec}} \neq \text{accepted}_i$ , set the 'lost' flag to  $\text{lost} \leftarrow \text{true}$ . This ensures that keys can only be tested once and, in case of internal keys, if they have just been accepted but not used yet, ensuring also that any partnered session that may have already established this key has not used it. If  $\text{label.rect\_auth}_i = \text{unauth}$ , or

---

<sup>7</sup>Note that for internal keys this implicitly assumes the following property of the later-defined Match security: Whenever two partnered sessions both accept a key in some stage, these keys will be equal.

if  $\text{label.rect\_auth}_i = \text{unilateral}$  and  $\text{label.role} = \text{responder}$ , but there is no session  $\text{label}'$  (for  $\text{label} \neq \text{label}'$ ) in  $\text{List}_5$  with  $\text{label.cid}_i = \text{label'.cid}_i$ , then set  $\text{lost} \leftarrow \text{true}$ . This ensures that having an honest contributive partner is a prerequisite for testing unauthenticated stages, resp. the responder sessions in a unilaterally authenticated stage.<sup>8</sup> The check is based on the uncorrupted authentication level  $\text{rect\_auth}_i$  in order to take corruptions between authentication upgrades into account.

Otherwise, set  $\text{label.tested}_i$  to true. If the test bit  $b_{\text{test}}$  is 0, sample a key  $K \leftarrow \mathcal{D}$  at random from the session key distribution  $\mathcal{D}$ . If  $b_{\text{test}} = 1$ , let  $K \leftarrow \text{label.key}_i$  be the real session key. If  $\text{USE}_i = \text{internal}$  (i.e., the tested  $i$ -th key is indicated as being used internally), set  $\text{label.key}_i \leftarrow K$ ; in other words, when  $b_{\text{test}} = 0$ , we replace an *internally* used session key by the random and independent test key  $K$  which is also used for consistent future usage *within* the key exchange protocol. In contrast, *externally used* session keys are not replaced by random ones, the adversary only receives the real (in case  $b_{\text{test}} = 1$ ) or random (in case  $b_{\text{test}} = 0$ ) key. This distinction between internal and external keys for **Test** queries emphasizes that external keys are not supposed to be used within the key exchange (and hence there is no need to register the tested random key in the protocol's session key field) while internal keys will be used (and hence the tested random key must be deployed in the remaining protocol steps for consistency).

Moreover, if there exists a partnered session  $\text{label}'$  which has also just accepted the  $i$ -th key (i.e.,  $\text{label.sid}_i = \text{label'.sid}_i$  and  $\text{label.st}_{\text{exec}} = \text{label'.st}_{\text{exec}} = \text{accepted}_i$ ), then also set  $\text{label'.tested}_i \leftarrow \text{true}$  and (only if  $\text{USE}_i = \text{internal}$ )  $\text{label'.key}_i \leftarrow \text{label.key}_i$  to ensure consistency (of later tests and (internal) key usage) in the special case that both  $\text{label}$  and  $\text{label}'$  are in state  $\text{accepted}_i$  and, hence, either of them can be tested first.

Return  $K$ .

### 4.3 Security of Multi-Stage Key Exchange Protocols

As in the formalization of the Bellare–Rogaway key exchange model by Brzuska et al. [BFWW11, Brz13], we model security according to two games, one for key indistinguishability, and one for session matching. The former is the classical notion of random-looking keys, refined under the term **Multi-Stage** security according to the advanced security aspects for multi-stage key exchange: (stage- $j$ ) forward secrecy, different authentication modes, and replayability. The **Match** property complements this notion by guaranteeing that the specified session identifiers  $\text{sid}$  effectively match the partnered sessions, and is likewise adapted to the multi-stage setting.

#### 4.3.1 Match Security

The notion of **Match** security ensures soundness of the session identifiers  $\text{sid}$ , i.e., that they properly identify partnered sessions in the sense that

1. sessions with the same session identifier for some stage hold the same key at that stage,
2. sessions with the same session identifier for some stage have opposite roles, except for potential multiple responders in replayable stages,
3. sessions with the same session identifier for some stage agree on that stage's authentication level,

---

<sup>8</sup>Note that  $\text{List}_5$  entries are only created for honest sessions, i.e., sessions generated by **NewSession** queries.

4. sessions with the same session identifier for some stage share the same contributive identifier at that stage,
5. sessions are partnered with the intended (authenticated) participant and, for mutual authentication based on pre-shared secrets, share the same key identifier,
6. session identifiers do not match across different stages, and
7. at most two sessions have the same session identifier at any non-replayable stage.

The Match security game  $G_{KE, \mathcal{A}}^{\text{Match}}$  thus is defined as follows.

**Definition 4.1** (Match security). *Let KE be a multi-stage key exchange protocol with properties (M, AUTH, FS, USE, REPLAY) and  $\mathcal{A}$  be a PPT adversary interacting with KE via the queries defined in Section 4.2 in the following game  $G_{KE, \mathcal{A}}^{\text{Match}}$ :*

**Setup.** *In the public-key variant (pMSKE), the challenger generates long-term public/private-key pairs for each participant  $U \in \mathcal{U}$ .*

**Query.** *The adversary  $\mathcal{A}$  receives the generated public keys (pMSKE) and has access to the queries NewSecret, NewSession, Send, Reveal, Corrupt, and Test.*

**Stop.** *At some point, the adversary stops with no output.*

*We say that  $\mathcal{A}$  wins the game, denoted by  $G_{KE, \mathcal{A}}^{\text{Match}} = 1$ , if at least one of the following conditions holds:*

1. *There exist two distinct labels label, label' such that  $\text{label.sid}_i = \text{label'.sid}_i \neq \perp$  for some stage  $i \in \{1, \dots, M\}$ , but  $\text{label.key}_i \neq \text{label'.key}_i$ . (Different session keys in some stage of partnered sessions.)*
2. *There exist two distinct labels label, label' such that  $\text{label.sid}_i = \text{label'.sid}_i \neq \perp$  for some stage  $i \in \{1, \dots, M\}$ , but  $\text{label.role} = \text{label'.role}$  and  $\text{REPLAY}_i = \text{nonreplayable}$ , or  $\text{label.role} = \text{label'.role} = \text{initiator}$  and  $\text{REPLAY}_i = \text{replayable}$ . (Non-opposite roles of partnered sessions in non-replayable stage.)*
3. *There exist two distinct labels label, label' such that  $\text{label.sid}_i = \text{label'.sid}_i \neq \perp$  for some stage  $i \in \{1, \dots, M\}$ , but  $\text{label.auth}_i \neq \text{label'.auth}_i$ . (Different authentication types in some stage of partnered sessions.)<sup>9</sup>*
4. *There exist two distinct labels label, label' such that  $\text{label.sid}_i = \text{label'.sid}_i \neq \perp$  for some stage  $i \in \{1, \dots, M\}$ , but  $\text{label.cid}_i \neq \text{label'.cid}_i$  or  $\text{label.cid}_i = \text{label'.cid}_i = \perp$ . (Different or unset contributive identifiers in some stage of partnered sessions.)*
5. *There exist two distinct labels label, label' such that  $\text{label.sid}_i = \text{label'.sid}_i \neq \perp$  and  $\text{label.sid}_j = \text{label'.sid}_j \neq \perp$  for stages  $i, j \in \{1, \dots, M\}$  where  $j \leq i$ , with  $\text{label.role} = \text{initiator}$  and  $\text{label'.role} = \text{responder}$  such that*
  - $\text{label.auth}_{j,1} \leq i$  (unilateral authentication), but  $\text{label.pid} \neq \text{label'.id}$ , or

---

<sup>9</sup>Observe that Match security ensures agreement on the *intended* authentication levels (including potential upgrades); the *rectified* authentication level in contrast is a technical element of the security model capturing the actual level achieved in light of early corruptions when evaluating Test queries.

- $\text{label.auth}_{j,2} \leq i$  (mutual authentication), but  $\text{label.id} \neq \text{label'.pid}$  or (only for *sMSKE*)  $\text{label.pssid} \neq \text{label'.pssid}$ .

(Different intended authenticated partner or (only *sMSKE*) different key identifiers in mutual authentication.)

6. There exist two (not necessarily distinct) labels  $\text{label}$ ,  $\text{label'}$  such that  $\text{label.sid}_i = \text{label'.sid}_j \neq \perp$  for some stages  $i, j \in \{1, \dots, M\}$  with  $i \neq j$ . (Different stages share the same session identifier.)
7. There exist three pairwise distinct labels  $\text{label}$ ,  $\text{label'}$ ,  $\text{label''}$  such that  $\text{label.sid}_i = \text{label'.sid}_i = \text{label''.sid}_i \neq \perp$  for some stage  $i \in \{1, \dots, M\}$  with  $\text{REPLAY}_i = \text{nonreplayable}$ . (More than two sessions share the same session identifier in a non-replayable stage.)

We say  $\text{KE}$  is Match-secure if for all PPT adversaries  $\mathcal{A}$  the following advantage function is negligible in the security parameter:

$$\text{Adv}_{\text{KE}, \mathcal{A}}^{\text{Match}} := \Pr \left[ G_{\text{KE}, \mathcal{A}}^{\text{Match}} = 1 \right].$$

### 4.3.2 Multi-Stage Security

The second and core notion, Multi-Stage security, captures Bellare–Rogaway-like key secrecy in the multi-stage setting as follows.

**Definition 4.2** (Multi-Stage security). Let  $\text{KE}$  be a multi-stage key exchange protocol with properties  $(M, \text{AUTH}, \text{FS}, \text{USE}, \text{REPLAY})$  and key distribution  $\mathcal{D}$ , and  $\mathcal{A}$  a PPT adversary interacting with  $\text{KE}$  via the queries defined in Section 4.2 in the following game  $G_{\text{KE}, \mathcal{A}}^{\text{Multi-Stage}, \mathcal{D}}$ :

**Setup.** The challenger chooses the test bit  $b_{\text{test}} \leftarrow_{\$} \{0, 1\}$  at random and sets  $\text{lost} \leftarrow \text{false}$ . In the public-key variant (*pMSKE*), it furthermore generates long-term public/private-key pairs for each participant  $U \in \mathcal{U}$ .

**Query.** The adversary  $\mathcal{A}$  receives the generated public keys (*pMSKE*) and has access to the queries `NewSecret`, `NewSession`, `Send`, `Reveal`, `Corrupt`, and `Test`. Recall that such queries may set  $\text{lost}$  to `true`.

**Guess.** At some point,  $\mathcal{A}$  stops and outputs a guess  $b$ .

**Finalize.** The challenger sets the ‘lost’ flag to  $\text{lost} \leftarrow \text{true}$  if there exist two (not necessarily distinct) labels  $\text{label}$ ,  $\text{label'}$  and some stage  $i \in \{1, \dots, M\}$  such that  $\text{label.sid}_i = \text{label'.sid}_i$ ,  $\text{label.st}_{\text{key}, i} = \text{revealed}$ , and  $\text{label'.tested}_i = \text{true}$ . (Adversary has tested and revealed the key of some stage in a single session or in two partnered sessions.)

We say that  $\mathcal{A}$  wins the game, denoted by  $G_{\text{KE}, \mathcal{A}}^{\text{Multi-Stage}, \mathcal{D}} = 1$ , if  $b = b_{\text{test}}$  and  $\text{lost} = \text{false}$ . Note that the winning condition is independent of forward secrecy and authentication properties of  $\text{KE}$ , as those are directly integrated in the affected (`Reveal` and `Corrupt`) queries and the finalization step of the game; for example, `Corrupt` is defined differently for non-forward-secrecy versus stage- $j$  forward secrecy.

We say  $\text{KE}$  is Multi-Stage-secure with properties  $(M, \text{AUTH}, \text{FS}, \text{USE}, \text{REPLAY})$  if  $\text{KE}$  is Match-secure and for all PPT adversaries  $\mathcal{A}$  the following advantage function is negligible in the security parameter:

$$\text{Adv}_{\text{KE}, \mathcal{A}}^{\text{Multi-Stage}, \mathcal{D}} := \Pr \left[ G_{\text{KE}, \mathcal{A}}^{\text{Multi-Stage}, \mathcal{D}} = 1 \right] - \frac{1}{2}.$$

## 5 Security Analysis of the TLS 1.3 Full 1-RTT Handshake

We now come to analyzing the TLS 1.3 full 1-RTT handshake in the public-key multi-stage key exchange (pMSKE) model.

**Protocol properties.** The full handshake targets the following protocol-specific properties (M, AUTH, FS, USE, REPLAY):

- **M = 6:** The full 1-RTT handshake consists of six stages deriving, in order: the client and server handshake traffic keys  $tk_{chs}$  and  $tk_{shs}$ , the client and server application traffic secrets CATS and SATS, the exporter master secret EMS, and the resumption master secret RMS. As shown in Figure 1, we consider all stages' keys being derived on either side as soon as the relevant main secret (ES, HS, MS) becomes available, despite client/server keys derived in parallel might become active with some delay based on the flow direction.
- **AUTH =  $\{((3, m), (3, m), (3, m), (4, m), (5, m), (6, m)) \mid m \in \{6, \infty\}\}$ :** The handshake traffic keys  $tk_{chs}/tk_{shs}$  are initially unauthenticated and all keys are unilaterally authenticated after stage 3 is reached. With (optional) client authentication, all keys furthermore become mutually authenticated with stage  $m = 6$ ; otherwise they never reach this level,  $m = \infty$ .
- **FS = 1:** The full 1-RTT handshake ensures forward secrecy for all keys derived.
- **USE = (internal :  $\{1, 2\}$ , external :  $\{3, 4, 5, 6\}$ ):** The handshake traffic keys are used internally to encrypt the second part of the handshake; all other keys are external.
- **REPLAY = (nonreplayable :  $\{1, 2, 3, 4, 5, 6\}$ ):** The keys of all stages are non-replayable in the full 1-RTT handshake.

**Session and contributive identifiers.** As part of the analysis in the pMSKE model, we need to define how session and contributive identifiers are set for each stage during execution of the TLS 1.3 full 1-RTT handshake.

Session identifiers are set upon acceptance of each stage and include a label and all handshake messages up to this point (entering the key derivation):

$$\begin{aligned}
 sid_1 &= (\text{"CHTS"}, CH, CKS, SH, SKS), \\
 sid_2 &= (\text{"SHTS"}, CH, CKS, SH, SKS), \\
 sid_3 &= (\text{"CATS"}, CH, CKS, SH, SKS, EE, CR^*, SCRT, SCV, SF), \\
 sid_4 &= (\text{"SATS"}, CH, CKS, SH, SKS, EE, CR^*, SCRT, SCV, SF), \\
 sid_5 &= (\text{"EMS"}, CH, CKS, SH, SKS, EE, CR^*, SCRT, SCV, SF), \\
 sid_6 &= (\text{"RMS"}, CH, CKS, SH, SKS, EE, CR^*, SCRT, SCV, SF, CCRT^*, CCV^*, CF).
 \end{aligned}$$

Here, starred (\*) components are present only in mutual authentication mode. Note that we define session identifiers over the *unencrypted* handshake messages.

For the contributive identifiers in stages 1 and 2, client (resp. server) upon sending (resp. receiving) the `ClientHello` and `ClientKeyShare` messages set  $cid_1 = (\text{"CHTS"}, CH, CKS)$ ,  $cid_2 = (\text{"SHTS"}, CH, CKS)$  and later, upon receiving (resp. sending) the `ServerHello` and `ServerKeyShare` messages, extend it to  $cid_1 = (\text{"CHTS"}, CH, CKS, SH, SKS)$ ,  $cid_2 = (\text{"SHTS"}, CH, CKS, SH, SKS)$ . All other contributive identifiers are set to  $cid_i = sid_i$  (for stages  $i \in \{3, 4, 5, 6\}$ ) when the respective session identifier is set.

## 5.1 Match Security

We are now ready to give our formal security results for the TLS 1.3 full 1-RTT handshake, beginning with Match security.

**Theorem 5.1** (Match security of TLS1.3-full-1RTT). *The TLS 1.3 full 1-RTT handshake is Match-secure with properties (M, AUTH, FS, USE, REPLAY) given above. For any efficient adversary  $\mathcal{A}$  we have*

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{\text{Match}} \leq n_s^2 \cdot \frac{1}{q} \cdot 2^{-|\text{nonce}|},$$

where  $n_s$  is the maximum number of sessions,  $q$  is the group order, and  $|\text{nonce}| = 256$  is the bit-length of the nonces.

Recall that Match security is a soundness property of the session identifiers. From our definition of session identifiers above, it follows immediately that partnered sessions agree on the derived key, opposite roles, authentication properties, contributive identifiers, and the respective stages. The security bound arises as the birthday bound for two honest sessions choosing the same nonce and group element; this not happening ensures at most two partners share the same session identifier.

*Proof.* We need to show the seven properties of Match security (cf. Definition 4.1).

1. *Sessions with the same session identifier for some stage hold the same key at that stage.*

The session identifiers in each stage include the Diffie–Hellman shares  $g^x$  and  $g^y$  (through the CKS and SKS messages, fixing the only key input  $\text{DHE} = g^{xy}$  to all derived stage keys (recall that  $\text{PSK} = 0$  in the TLS 1.3 full 1-RTT handshake). Furthermore, for each stage, the session identifier includes all handshake messages that enter the key derivation: for stages 1 and 2 messages up to SKS, for stages 3–5 messages up to SF, and for stage 6 all messages (up to CF). In each stage, the session identifier hence determines *all* inputs to the key derivation, and agreement on it thus ensures agreement on the stage key.

2. *Sessions with the same session identifier for some stage have opposite roles, except for potential multiple responders in replayable stages.*

Assuming at most two sessions share the same session identifier (which we show below), two initiator (client) or responder (server) sessions never hold the same session identifier as they never accept wrong-role incoming messages, and the initial Hello messages are typed with the sender’s role. There are no replayable stages in the TLS 1.3 full 1-RTT handshake.

3. *Sessions with the same session identifier for some stage agree on that stage’s authentication level.*

By definition, the authentication for stages 1–2 and 3–5 are fixed to `unauth` and `unilateral` (from stage 3 on), respectively, hence agreed upon by all sessions. For the last stage, the presence of CR, CCRT, and CCV in `sid6` unambiguously determines if, from stage 6 on, keys are mutually authenticated (and unilaterally otherwise).

4. *Sessions with the same session identifier for some stage share the same contributive identifier.*

This holds due to, for each stage  $i$ , the contributive identifier `cidi` being final and equal to `sidi` once the session identifier is set.

5. *Sessions are partnered with the intended (authenticated) participant.*

This case only applies to unilaterally or mutually authenticated stages, which is achieved, possibly retroactively, when reaching stages 3, resp. stage 6 (only if the client authenticates).



In the TLS 1.3 full 1-RTT handshake, peer identities are learned through the **Certificate** messages. As we are only concerned with honest client and server sessions for **Match** security, which will only send certificates attesting their own identity, agreement on **SCRT** ensures agreeing on the server (responder) identity, and vice versa for **CCRT** and the client (initiator) identity. Such agreement is ensured through including **SCRT** in the session identifier for stage 3 for unilateral authentication, and **SCRT** and **CCRT** for mutual authentication in  $\text{sid}_6$ : once two sessions reach these stages and agree on  $\text{sid}_3$ , resp.  $\text{sid}_6$ , they (retroactively) also agree on the intended (responder, resp. initiator) peer.

6. *Session identifiers are distinct for different stages.*

This holds trivially as each stage's session identifier has a unique label.

7. *At most two sessions have the same session identifier at any non-replayable stage.*

Recall that all session identifiers held by some session include that session's random nonce and Diffie–Hellman share. Therefore, for a threefold collision among session identifiers of honest parties, some session would need to pick the same group element and nonce as one other session (which then may be partnered through a regular protocol run to some third session). The probability for such collision to happen can be bounded from above by the birthday bound  $n_s^2 \cdot 1/q \cdot 2^{-|\text{nonce}|}$ , where  $n_s$  is the maximum number of sessions,  $q$  is the group order, and  $|\text{nonce}| = 256$  the nonces' bit-length.  $\square$

## 5.2 Multi-Stage Security

We now come to the core multi-stage security result for the TLS 1.3 full 1-RTT handshake.

**Theorem 5.2** (Multi-Stage security of TLS1.3-full-1RTT). *The TLS 1.3 full 1-RTT handshake is Multi-Stage-secure with properties (M, AUTH, FS, USE, REPLAY) given above. Formally, for any efficient adversary  $\mathcal{A}$  against the Multi-Stage security there exist efficient algorithms  $\mathcal{B}_1, \dots, \mathcal{B}_7$  such that*

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{\text{Multi-Stage}, \mathcal{D}} \leq 6n_s \left( \begin{aligned} &\text{Adv}_{\mathcal{H}, \mathcal{B}_1}^{\text{COLL}} + n_u \cdot \text{Adv}_{\text{SIG}, \mathcal{B}_2}^{\text{EUF-CMA}} \\ &+ n_s \left( \begin{aligned} &\text{Adv}_{\text{HKDF.Extract}, \mathcal{G}, \mathcal{B}_3}^{\text{dual-snPRF-ODH}} + \text{Adv}_{\text{HKDF.Expand}, \mathcal{B}_4}^{\text{PRF-sec}} \\ &+ 2 \cdot \text{Adv}_{\text{HKDF.Expand}, \mathcal{B}_5}^{\text{PRF-sec}} + \text{Adv}_{\text{HKDF.Extract}, \mathcal{B}_6}^{\text{PRF-sec}} \\ &+ \text{Adv}_{\text{HKDF.Expand}, \mathcal{B}_7}^{\text{PRF-sec}} \end{aligned} \right) \end{aligned} \right)$$

where  $n_s$  is the maximum number of sessions and  $n_u$  is the maximum number of users.

For the TLS 1.3 full 1-RTT handshake, Multi-Stage security essentially follows from two lines of reasoning. First, the (unforgeable) signatures covering (a collision-resistant hash of) the full **Hello** messages ensure that session stages with an authenticated peer share exchanged Diffie–Hellman values originating from an honest partner session. Then, all keys are derived in a way ensuring that (a) from a Diffie–Hellman secret unknown to the adversary sessions derive keys indistinguishable from random (under PRF-ODH and PRF assumptions on the HKDF.Extract and HKDF.Expand steps) which (b) are independent, allowing revealing and testing of session keys across different stages.

*Proof.* In the following, we proceed via a sequence of games. Starting from the Multi-Stage game, we bound the advantage difference of adversary  $\mathcal{A}$  between any two games by complexity-theoretic assumptions until we reach a game where the adversary  $\mathcal{A}$  cannot win, i.e., its advantage is at most 0.

**Game 0.** This is the original Multi-Stage game, i.e.,

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{\text{Multi-Stage}, \mathcal{D}} = \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_0}.$$

**Game 1.** In a first step, we restrict the adversary  $\mathcal{A}$  in the Multi-Stage game to make only a single Test query. That is we can formally turn any multi-query adversary  $\mathcal{A}$  into an adversary  $\mathcal{A}_1$  which makes only a single Test query. This reduces its advantage, based on a careful hybrid argument, by a factor at most  $1/6n_s$  for the six stages in each of the  $n_s$  sessions. Note that in the hybrid argument  $\mathcal{A}_1$  randomly guesses one of the sessions in advance and only performs the single Test query for this session. The other Test queries of a multi-query attacker are gradually substituted by carefully crafted Reveal queries, where the single-query attacker  $\mathcal{A}_1$  needs to know the correct partnering of sessions via session identifiers  $\text{sid}$  for a correct simulation, e.g., to avoid losses due to bad Reveal-Test combinations on session partners due to the new Reveal queries. The session identifiers  $\text{sid}_1$  and  $\text{sid}_2$  only contain public information such that partnering is easy to check for them. But then handshake encryption is turned on such that  $\text{sid}_3, \dots, \text{sid}_6$  are based on encrypted data. Fortunately, if the single Test query concerns a (client or server) handshake traffic secret then partnering is easy to decide based on  $\text{sid}_1$  resp.  $\text{sid}_2$ . If the Test query refers to a later key we can reveal the handshake traffic keys of earlier stages, use them to decrypt the subsequent communication, and hence determine  $\text{sid}_3, \dots, \text{sid}_6$  as well. We provide the full details of this hybrid argument in Appendix A.

Incorporating the transformation of  $\mathcal{A}$  into  $\mathcal{A}_1$  into the game, i.e., by having the challenger guess the right session and making the adaptations, we get

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_0} \leq 6n_s \cdot \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_1}.$$

From now on, we can refer to *the* session label tested at stage  $i$ , and we can assume that we know this session number (according to the order of initiated sessions) at the outset of the experiment.

**Game 2.** In this game, the challenger aborts if any two honest sessions compute the same hash value for different inputs in any evaluation of the hash function  $H$ . We can break the collision-resistance of  $H$  in case of this event by letting a reduction  $\mathcal{B}_1$  output the two distinct input values to  $H$ . Thus:

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_1} \leq \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_2} + \text{Adv}_{H, \mathcal{B}_1}^{\text{COLL}}.$$

From here on, our security analysis separately considers the two (disjoint) cases that

- A. the tested session label has no honest contributive partner in the first stage (i.e., there exists no  $\text{label}' \neq \text{label}$  with  $\text{label}.cid_1 = \text{label}'.cid_1$ ), and
- B. the tested session label has an honest contributive partner in the first stage (i.e., there exists  $\text{label}'$  with  $\text{label}.cid_1 = \text{label}'.cid_1$ ).

This allows us to consider the adversary's advantage separately for these two cases A (denoted “test w/o partner”) and B (“test w/ partner”):

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_2} \leq \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_2, \text{ test w/o partner}} + \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_2, \text{ test w/ partner}}.$$

### Case A. Test without Partner

We first consider the case that the tested session `label` has no stage-1 contributive partner, which implies it does not have a contributive partner in any stage. By definition, an adversary cannot win if the `Test` query issued to such session is in a stage that, at the time of the test query, has an unauthenticated peer. Here, authentication refers to the *rectified* level, because the `Test` oracle checks against this refined property. Hence, for a tested client session, `Test` (for any stage) cannot be issued before stage 3 is reached and later only if corruption of the client or the partnered server has not taken place before stage 3. Else the adversary loses the game. For a server session, `Test` can only be issued when stage 6 is reached and client authentication is performed. Here, again, the client cannot be corrupted earlier, else the rectified authentication level would be unauthenticated.

**Game A.0.** Equals  $G_2$  with adversary restricted to test a session without honest contributive partner in the first stage.

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_2, \text{ test w/o partner}} = \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{A.0}}$$

**Game A.1.** In this game, we let the challenger guess the peer identity  $U \in \mathcal{U}$  of the tested session `label` (observe that one must be set in order for `Test` to be admissible, as discussed above), and abort if that guess was incorrect (i.e., `label.pid`  $\neq U$ ). This can reduce  $\mathcal{A}$ 's advantage by a factor at most the number of users  $n_u$ :

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{A.0}} \leq n_u \cdot \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{A.1}}$$

**Game A.2.** We now let the challenger abort the game if the tested session `label` receives, within the `CertificateVerify` message from its peer `label.pid = U`, a valid signature on some (hash value of a) message that has not been computed by any honest session of user  $U$ . Note that this message must include the transcript data `ClientHello` || ... || `ClientCert` resp. `ClientHello` || ... || `ServerCert` (cf. Table 2). Observe that, as discussed above, when the `Test` query is issued to `label`, such a message must have been received, in the case of a client, prior to accepting stage 3 and with no previous corruption of the server; or, in the case of a server, prior to stage 6 when the server is talking to an authenticating client which is not corrupted yet.

We can bound the probability of Game  $G_{A.2}$  aborting for this reason by the advantage of an adversary  $\mathcal{B}_2$  against the EUF-CMA security of the signature scheme `SIG`. In the reduction  $\mathcal{B}_2$  receives a public key  $pk_U$  of a signature scheme, computes the long-term keys of all parties  $U' \in \mathcal{U} \setminus \{U\}$  except  $U$  and simulates  $G_{A.1}$  for  $\mathcal{A}_1$ . Whenever in that simulation  $\mathcal{B}_2$  has to compute a signature under  $sk_U$ , it does so via its signing oracle. When `label` receives a valid signature  $\sigma$  on the (hash value of the) message  $m$ , adversary  $\mathcal{B}_2$  outputs  $(H(m), \sigma)$  as its forgery. Note that at this point the partnered session cannot be corrupted such that the signature forger does not need to reveal the secret signing key before outputting the forgery.

It remains to argue that the pair  $(H(m), \sigma)$  constitutes a successful forgery. To see this note that the tested session `label` computes the hash value  $H(m)$  of the message  $m$  to verify correctness, but such that no other honest session has computed a signature for this message. According to Game  $G_2$ , this also means that no other honest session has derived the same hash value  $H(m') = H(m)$  for some other message  $m'$ . We conclude that the hash value  $H(m)$  has not been signed by user  $U$  before.

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{A.1}} \leq \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{A.2}} + \text{Adv}_{\text{SIG}, \mathcal{B}_2}^{\text{EUF-CMA}}$$

It follows for Case A that the adversary cannot make a legitimate **Test** query at all, unless it forges signatures. Either the sessions do not have a contributive partner, or the sessions in later stages have rejected because of invalid signatures. If the adversary cannot test any session without a contributive partner, it clearly has no advantage in predicting the secret challenge bit  $b$ :

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{A,2}} = 0.$$

## Case B. Test with Partner

**Game B.0.** This is  $G_2$  where the adversary is restricted to issuing a **Test** query to a session with an honest contributive partner in the first stage.

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_2, \text{ test w/ partner}} = \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B,0}}.$$

**Game B.1.** In this game, we guess a session  $\text{label}' \neq \text{label}$  (from at most  $n_s$  sessions in the game) and abort the game if  $\text{label.cid}_1 \neq \text{label'.cid}_1$  i.e. that  $\text{label}'$  is not the honest contributive partner in stage 1 of the tested session (recall that we assume such partner exists in this proof case). This reduces the adversary's advantage by a factor of at most  $1/n_s$ .

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B,0}} \leq n_s \cdot \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B,1}}.$$

**Game B.2.** In this game, we replace the handshake secret  $\text{HS}$  derived in the tested session and its contributive partner session with a uniformly random and independent string  $\widetilde{\text{HS}} \leftarrow \{0, 1\}^\lambda$ . We employ the **dual-snPRF-ODH** assumption (Definition 2.3) in order to be able to simulate the computation of  $\text{HS}$  in a partnered client session for a modified **ServerKeyShare** message. More precisely, we can turn any adversary capable of distinguishing this change into an adversary  $\mathcal{B}_3$  against the **dual-snPRF-ODH** security of the  $\text{HKDF.Extract}$  function (taking  $\text{dES}$  as first and  $\text{DHE}$  as second input). For this  $\mathcal{B}_3$  asks for a PRF challenge on  $\text{dES}$  computed in the test session and its honest contributive partner. It uses the obtained Diffie-Hellman shares  $g^x, g^y$  within **ClientKeyShare** and **ServerKeyShare** of the tested and contributive sessions, and the PRF challenge value as  $\text{HS}$  in the tested session. If necessary,  $\mathcal{B}_3$  uses its PRF-ODH queries to derive  $\text{HS}$  in the partnered session on differing  $g^{y'} \neq g^y$ . Providing a sound simulation of either  $G_{B,1}$  (if the bit sampled by the **dual-snPRF-ODH** challenger was 0 and thus  $\widetilde{\text{HS}} = \text{HKDF.Extract}(\text{dES}, g^{xy})$ ), or  $G_{B,2}$  (if the bit sampled by the **dual-snPRF-ODH** challenger was 1 and thus  $\widetilde{\text{HS}} \leftarrow \{0, 1\}^\lambda$ ), this bounds the advantage difference of  $\mathcal{A}$  as:

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B,1}} \leq \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B,2}} + \text{Adv}_{\text{HKDF.Extract}, \mathcal{G}, \mathcal{B}_3}^{\text{dual-snPRF-ODH}}.$$

**Game B.3.** In this game, we replace the pseudorandom function  $\text{HKDF.Expand}$  in all evaluations using the value  $\widetilde{\text{HS}}$  replaced in  $G_{B,2}$ . This affects the derivation of the client handshake traffic secret  $\text{CHTS}$ , the server handshake traffic secret  $\text{SHTS}$  and the derived handshake secret  $\text{dHS}$  in the target session and its matching partner, and the derived handshake secret  $\text{dHS}$  in all sessions using the same handshake secret  $\text{HS}$ . Note that for  $\text{CHTS}$  and  $\text{SHTS}$ , these values are distinct from any other session using the same handshake secret value  $\widetilde{\text{HS}}$ , as the evaluation also takes as input the hash value  $H_2 = \text{H}(\text{CH} \parallel \text{SH})$ , (where  $\text{CH}$  and  $\text{SH}$  contain the client and server random values  $r_c, r_s$  respectively) and by Game  $G_2$  we exclude hash collisions. We replace the derivation

of CHTS, SHTS and dHS in such sessions with random values  $\widetilde{\text{CHTS}}, \widetilde{\text{SHTS}}, \widetilde{\text{dHS}} \leftarrow_{\$} \{0, 1\}^\lambda$ . To ensure consistency, we replace derivations of dHS with the replaced  $\widetilde{\text{dHS}}$  sampled by the first session to evaluate HKDF.Expand using  $\widetilde{\text{HS}}$ . We can bound the difference that this step introduces in the advantage of  $\mathcal{A}$  by the security of the pseudorandom function HKDF.Expand. Note that by the previous game,  $\widetilde{\text{HS}}$  is a uniformly random value, and the replacement is sound. Thus:

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B.2}} \leq \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B.3}} + \text{Adv}_{\text{HKDF.Expand}, \mathcal{B}_4}^{\text{PRF-sec}}.$$

At this point,  $\widetilde{\text{CHTS}}$  and  $\widetilde{\text{SHTS}}$  are independent of any values computed in any session non-partnered (in stage 1 or 2) with the tested session: distinct session identifiers and no hash collisions (as of Game  $G_2$ ) ensure that the PRF label inputs for deriving  $\widetilde{\text{CHTS}}$  and  $\widetilde{\text{SHTS}}$  are unique.

**Game B.4.** In this game, we replace the pseudorandom function HKDF.Expand in all evaluations using the values  $\widetilde{\text{CHTS}}, \widetilde{\text{SHTS}}$  replaced in  $G_{B.3}$ . This affects the derivation of the client handshake traffic key  $\text{tk}_{\text{chs}}$ , and the server handshake traffic key  $\text{tk}_{\text{shs}}$  in the target session and its contributive partner. We replace the derivation of  $\text{tk}_{\text{chs}}$  and  $\text{tk}_{\text{shs}}$  with random values  $\widetilde{\text{tk}}_{\text{chs}} \leftarrow_{\$} \{0, 1\}^L$  and  $\widetilde{\text{tk}}_{\text{shs}} \leftarrow_{\$} \{0, 1\}^L$ , where  $L$  indicates the sum of key length and iv length for the negotiated AEAD scheme. We can bound the difference that this step introduces in the advantage of  $\mathcal{A}$  by the security of two evaluations of the pseudorandom functions HKDF.Expand. Note that by the previous game  $\widetilde{\text{CHTS}}$  and  $\widetilde{\text{SHTS}}$  are uniformly random values, and these replacements are sound. Thus:

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B.3}} \leq \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B.4}} + 2 \cdot \text{Adv}_{\text{HKDF.Expand}, \mathcal{B}_5}^{\text{PRF-sec}}.$$

**Game B.5.** In this game, we replace the pseudorandom function HKDF.Extract in all evaluations of the value  $\widetilde{\text{dHS}}$  replaced in  $G_{B.3}$ . This affects the derivation of the master secret MS in any session using the same derived handshake secret  $\widetilde{\text{dHS}}$ . We replace the derivation of MS in such sessions with the random value  $\widetilde{\text{MS}} \leftarrow_{\$} \{0, 1\}^\lambda$ . We can bound the difference that this step introduces in the advantage of  $\mathcal{A}$  by the security of the pseudorandom function HKDF.Extract. Note that by  $G_{B.3}$ ,  $\widetilde{\text{dHS}}$  is a uniformly random value and this replacement is sound. Thus:

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B.4}} \leq \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B.5}} + \text{Adv}_{\text{HKDF.Extract}, \mathcal{B}_6}^{\text{PRF-sec}}.$$

**Game B.6.** In this game, we replace the pseudorandom function HKDF.Expand in all evaluations of the value  $\widetilde{\text{MS}}$  replaced in  $G_{B.5}$  in the targeted session and its matching session. This affects the derivation of the client application traffic secret CATS, the server application traffic secret SATS, the exporter master secret EMS and the resumption master secret RMS. For CATS, SATS and EMS, these evaluations are distinct from any session non-partnered with the tested session, as the evaluation of HKDF.Expand also takes as input  $H_4 = \text{H}(\text{CH} \parallel \dots \parallel \text{SF})$  (where CH and SH contain the client and server random values  $r_c$  and  $r_s$  respectively), and by Game  $G_2$  we exclude hash collisions. For RMS, this evaluation is distinct from any session non-partnered with the tested session, as the evaluation of HKDF.Expand also takes as input  $H_5 = \text{H}(\text{CH} \parallel \dots \parallel \text{CF})$ . We replace the derivation of CATS, SATS, EMS and RMS with random values  $\widetilde{\text{CATS}}, \widetilde{\text{SATS}}, \widetilde{\text{EMS}}, \widetilde{\text{RMS}} \leftarrow_{\$} \{0, 1\}^\lambda$ . We can bound the difference that this step introduces in the advantage of  $\mathcal{A}$  by the secret of the pseudorandom function HKDF.Expand. Note that by the previous game  $\widetilde{\text{MS}}$  is a uniformly random and independent value, and these replacements are sound. Thus:

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B.5}} \leq \text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B.6}} + \text{Adv}_{\text{HKDF.Expand}, \mathcal{B}_7}^{\text{PRF-sec}}.$$

We note that in this game we have now replaced all stages' keys in the tested session with uniformly random values which, in the protocol execution, are independent of values in any non-partnered session to the tested session. Thus:

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_{B.6}} = 0.$$

Combining the given single bounds yields the security statement below:

$$\text{Adv}_{\text{TLS1.3-full-1RTT}, \mathcal{A}}^{G_2, \text{ test w/ partner}} \leq n_s \left( \begin{aligned} &\text{Adv}_{\text{HKDF.Extract}, \mathcal{G}, \mathcal{B}_3}^{\text{dual-snPRF-ODH}} + \text{Adv}_{\text{HKDF.Expand}, \mathcal{B}_4}^{\text{PRF-sec}} + 2 \cdot \text{Adv}_{\text{HKDF.Expand}, \mathcal{B}_5}^{\text{PRF-sec}} \\ &+ \text{Adv}_{\text{HKDF.Extract}, \mathcal{B}_6}^{\text{PRF-sec}} + \text{Adv}_{\text{HKDF.Expand}, \mathcal{B}_7}^{\text{PRF-sec}} \end{aligned} \right) \square$$

## 6 Security Analysis of the TLS 1.3 PSK/PSK-(EC)DHE (with Optional 0-RTT) Handshakes

We now turn to analyzing the TLS 1.3 pre-shared key handshakes, with and without Diffie–Hellman key exchange (PSK-(EC)DHE, resp. PSK) and with optional 0-RTT keys, in the pre-shared-secret multi-stage key exchange (sMSKE) model.

**Protocol properties.** The PSK/PSK-(EC)DHE (0-RTT) handshakes targets the following protocol-specific properties (M, AUTH, FS, USE, REPLAY):

- M = 8: The PSK handshakes with optional 0-RTT consist of eight stages deriving, in order: the early traffic secret ETS and early exporter master secret EEMS (both only in 0-RTT mode), the client and server handshake traffic keys  $\text{tk}_{\text{chs}}$  and  $\text{tk}_{\text{shs}}$ , the client and server application traffic secrets CATS and SATS, the exporter master secret EMS, and the resumption master secret RMS.
- The authentication properties AUTH differ between the PSK(-only) and the PSK-(EC)DHE (0-RTT) handshakes:
  - for PSK (0-RTT),  $\text{AUTH} = \{((1, 1), (2, 2), \dots, (8, 8))\}$ : All keys are immediately mutually authenticated (from the preshared key).
  - for PSK-(EC)DHE (0-RTT),  $\text{AUTH} = \{((1, 1), (2, 2), (5, 8), (5, 8), (5, 8), (6, 8), (7, 8), (8, 8))\}$ : The 0-RTT keys ETS/EEMS are always mutually authenticated, the handshake traffic keys  $\text{tk}_{\text{chs}}/\text{tk}_{\text{shs}}$  are initially unauthenticated, all non-0-RTT keys reach unilateral authentication with stage 5 and mutual authentication with stage 8.<sup>10</sup>
- Forward secrecy of the PSK handshake depends on whether an ephemeral Diffie–Hellman key exchange is performed:
  - for PSK-only,  $\text{FS} = \infty$ : The PSK-only handshake does not provide any forward secrecy.
  - for PSK-(EC)DHE,  $\text{FS} = 3$ : The PSK-(EC)DHE handshake provides forward secrecy for all non-0-RTT keys.

---

<sup>10</sup>It is not straightforward to see why some PSK-(EC)DHE keys are not considered to be immediately mutually authenticated, in contrast to keys from the PSK-only handshake. Consider the handshake traffic keys in the PSK-(EC)DHE handshake: in the model, the adversary  $\mathcal{A}$  could send its own  $g^x$  share to a server session; the server will derive the handshake traffic keys from PSK and DHE. Those keys should now be considered forward secret (due to the ephemeral DH shares), however when  $\mathcal{A}$  corrupts PSK, it can compute the handshake traffic keys. Hence, these keys cannot be treated as forward secret and mutually authenticated at the same time.